

Bitmap Indexes

Marko Kevac
Gophercon Russia 2019

<http://bit.ly/bitmapindexes>
<https://github.com/mkevac/gopherconrussia2019>

Contents

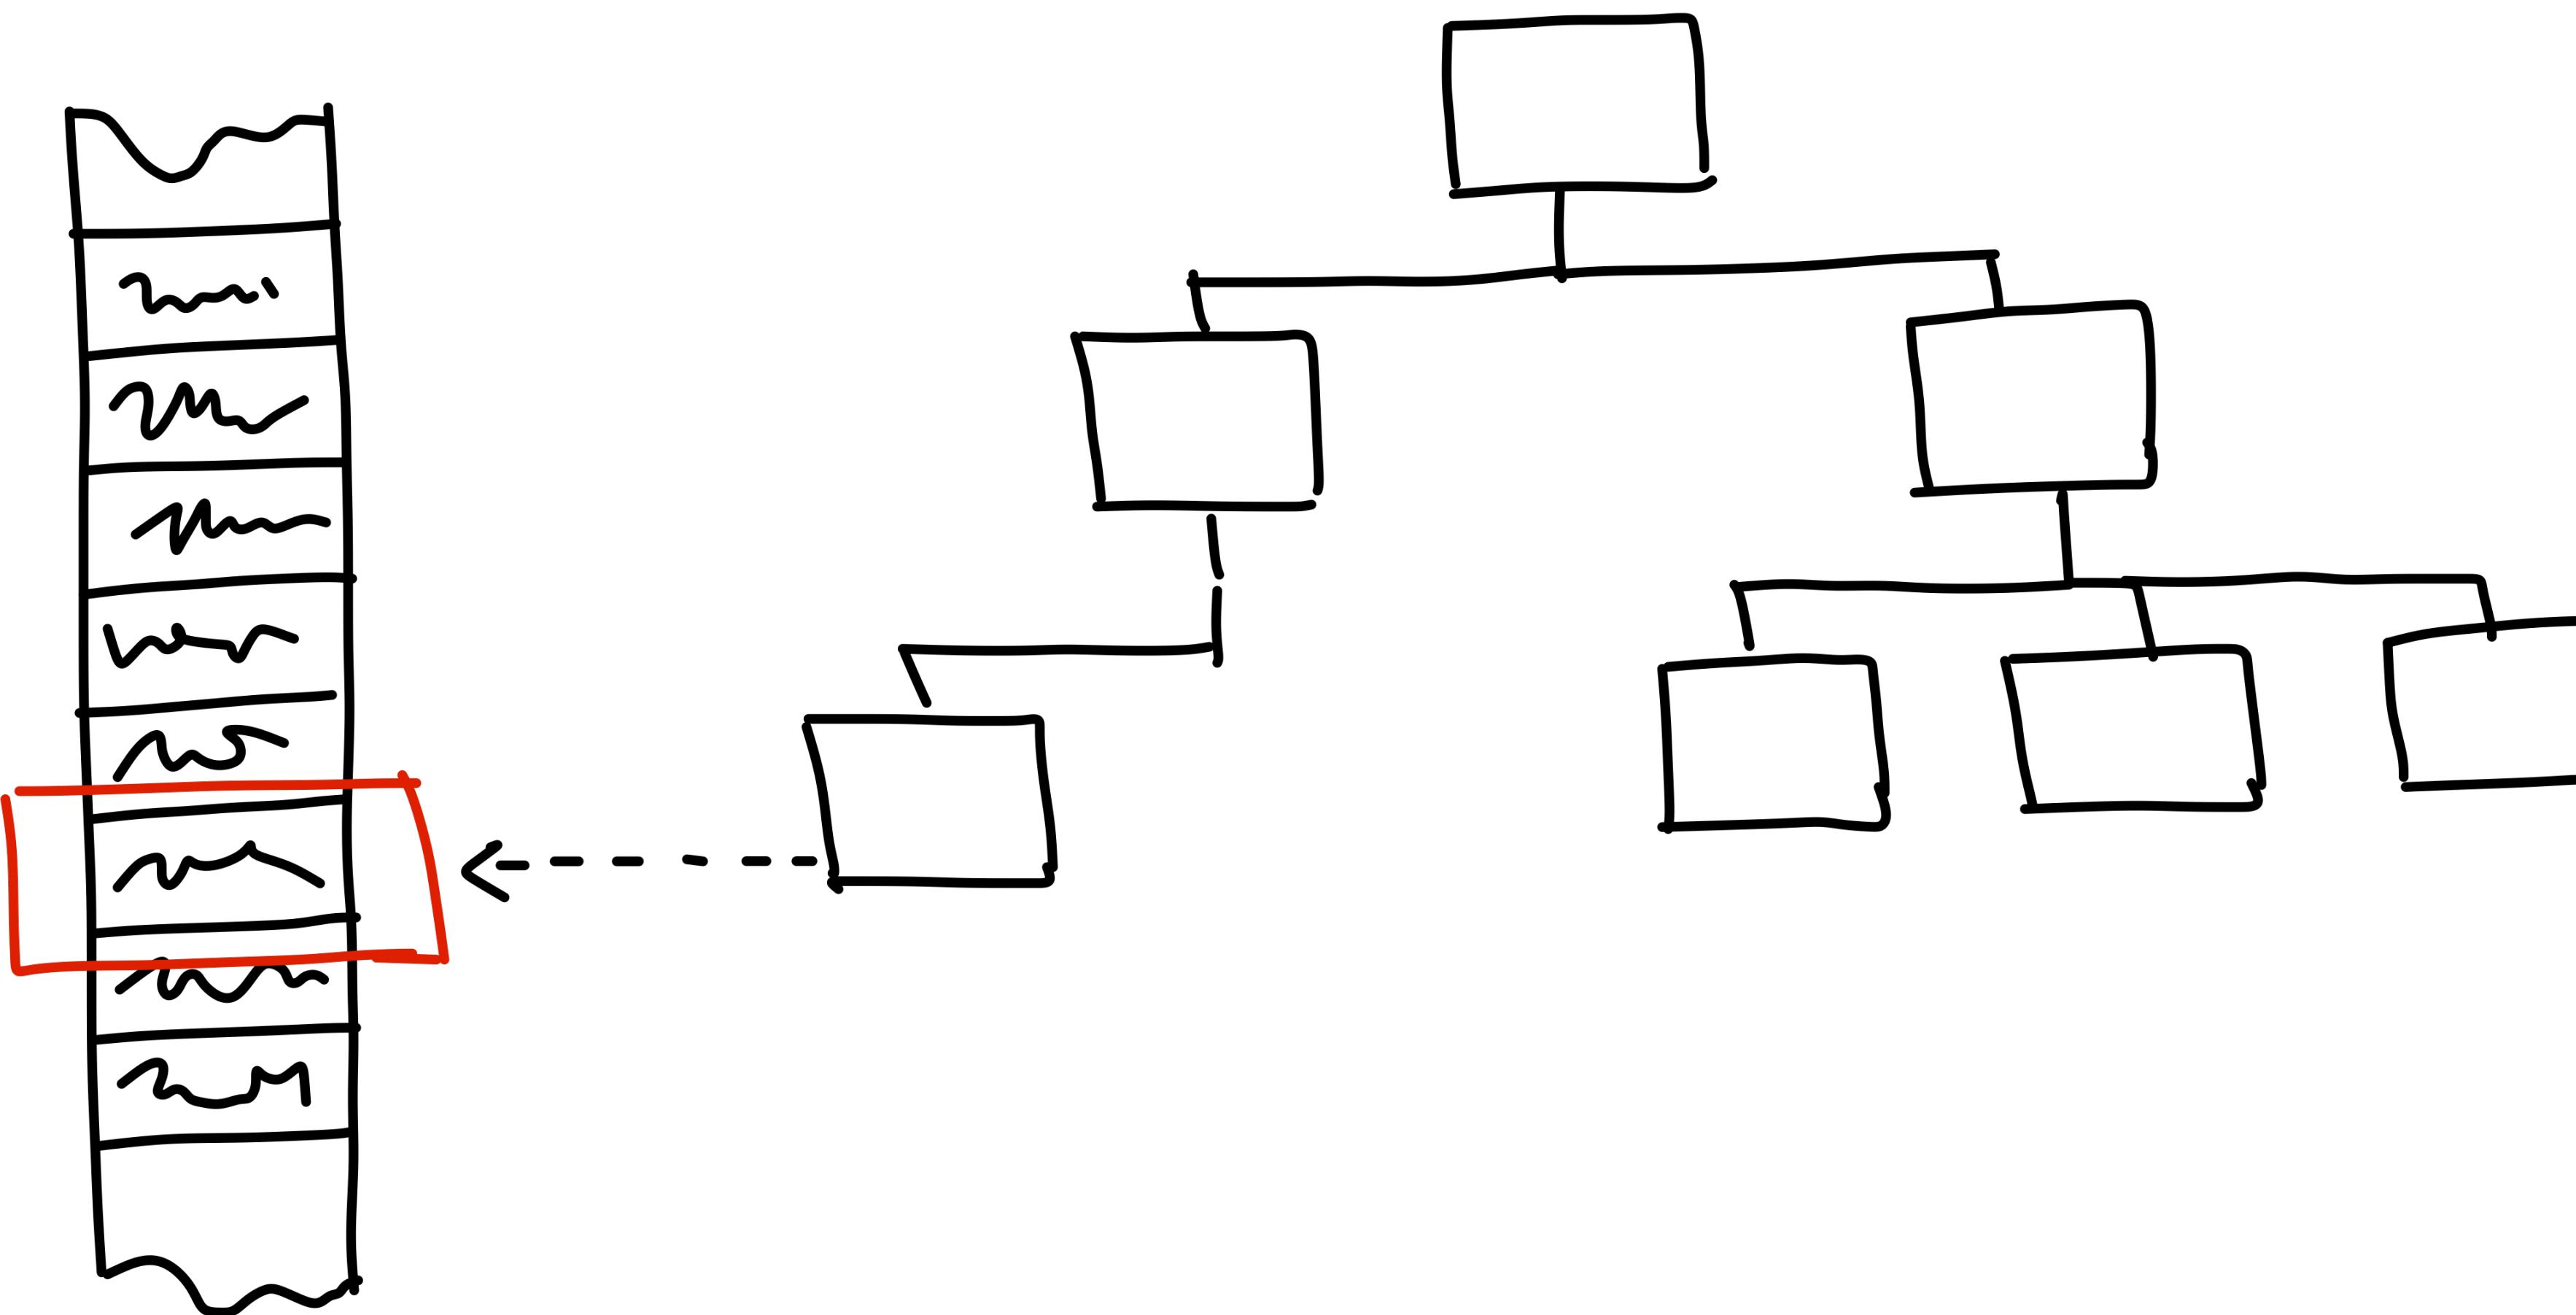
- 1.What indexes are.
- 2.What is a bitmap index.
- 3.Where it's used. Why it's not.
- 4.Simple implementation in Go. Fight with compiler.
- 5.Simple implementation in Go assembly.
- 6.Addressing the “problems”. *Also fun!*
- 7.Existing solutions.

FUN!

MEH..

*BORING
THEORY*

Indexes



Indexing approaches

Hierarchical division

- *-trees
- B-trees
- R-trees

$O(\log N)$



Indexing approaches

Hash mapping

Hash maps
Reverse indexes

O(1)





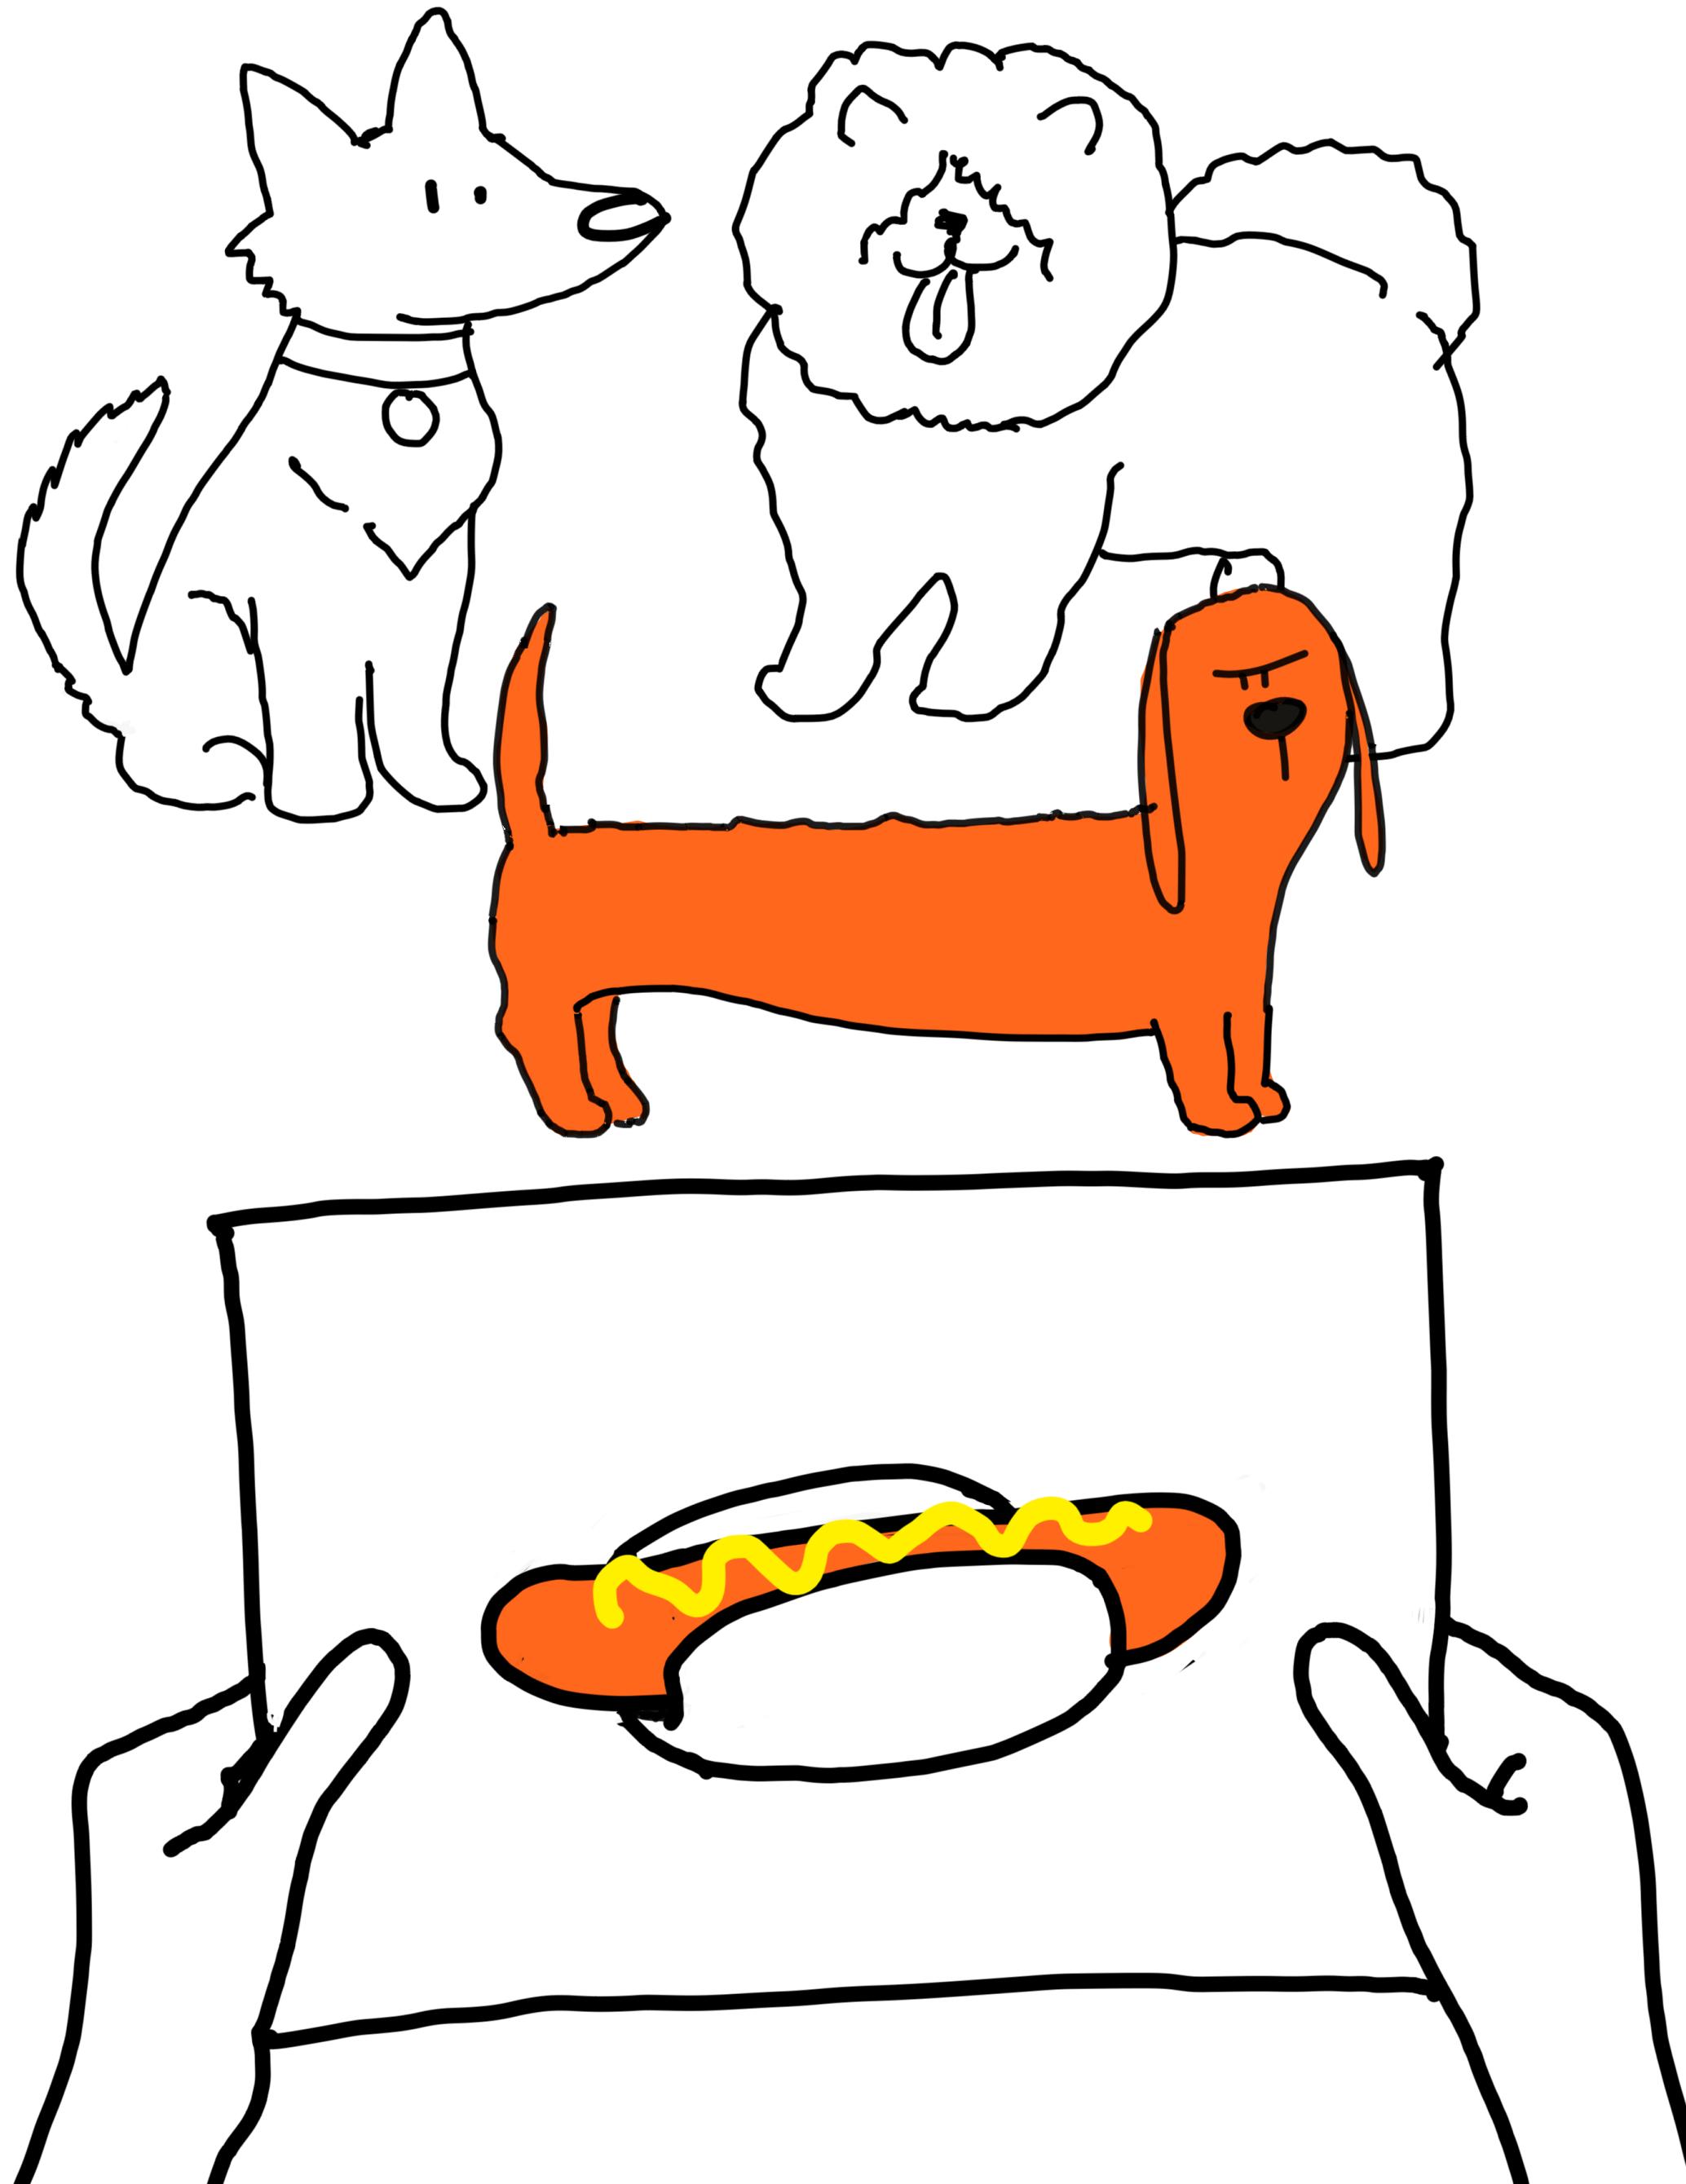
Indexing approaches

Instantly knowing

Bloom filters

Cuckoo filters

$O(1)$



Indexing approaches

Fully using hardware capabilities

Bitmap indexes

$O(N)$

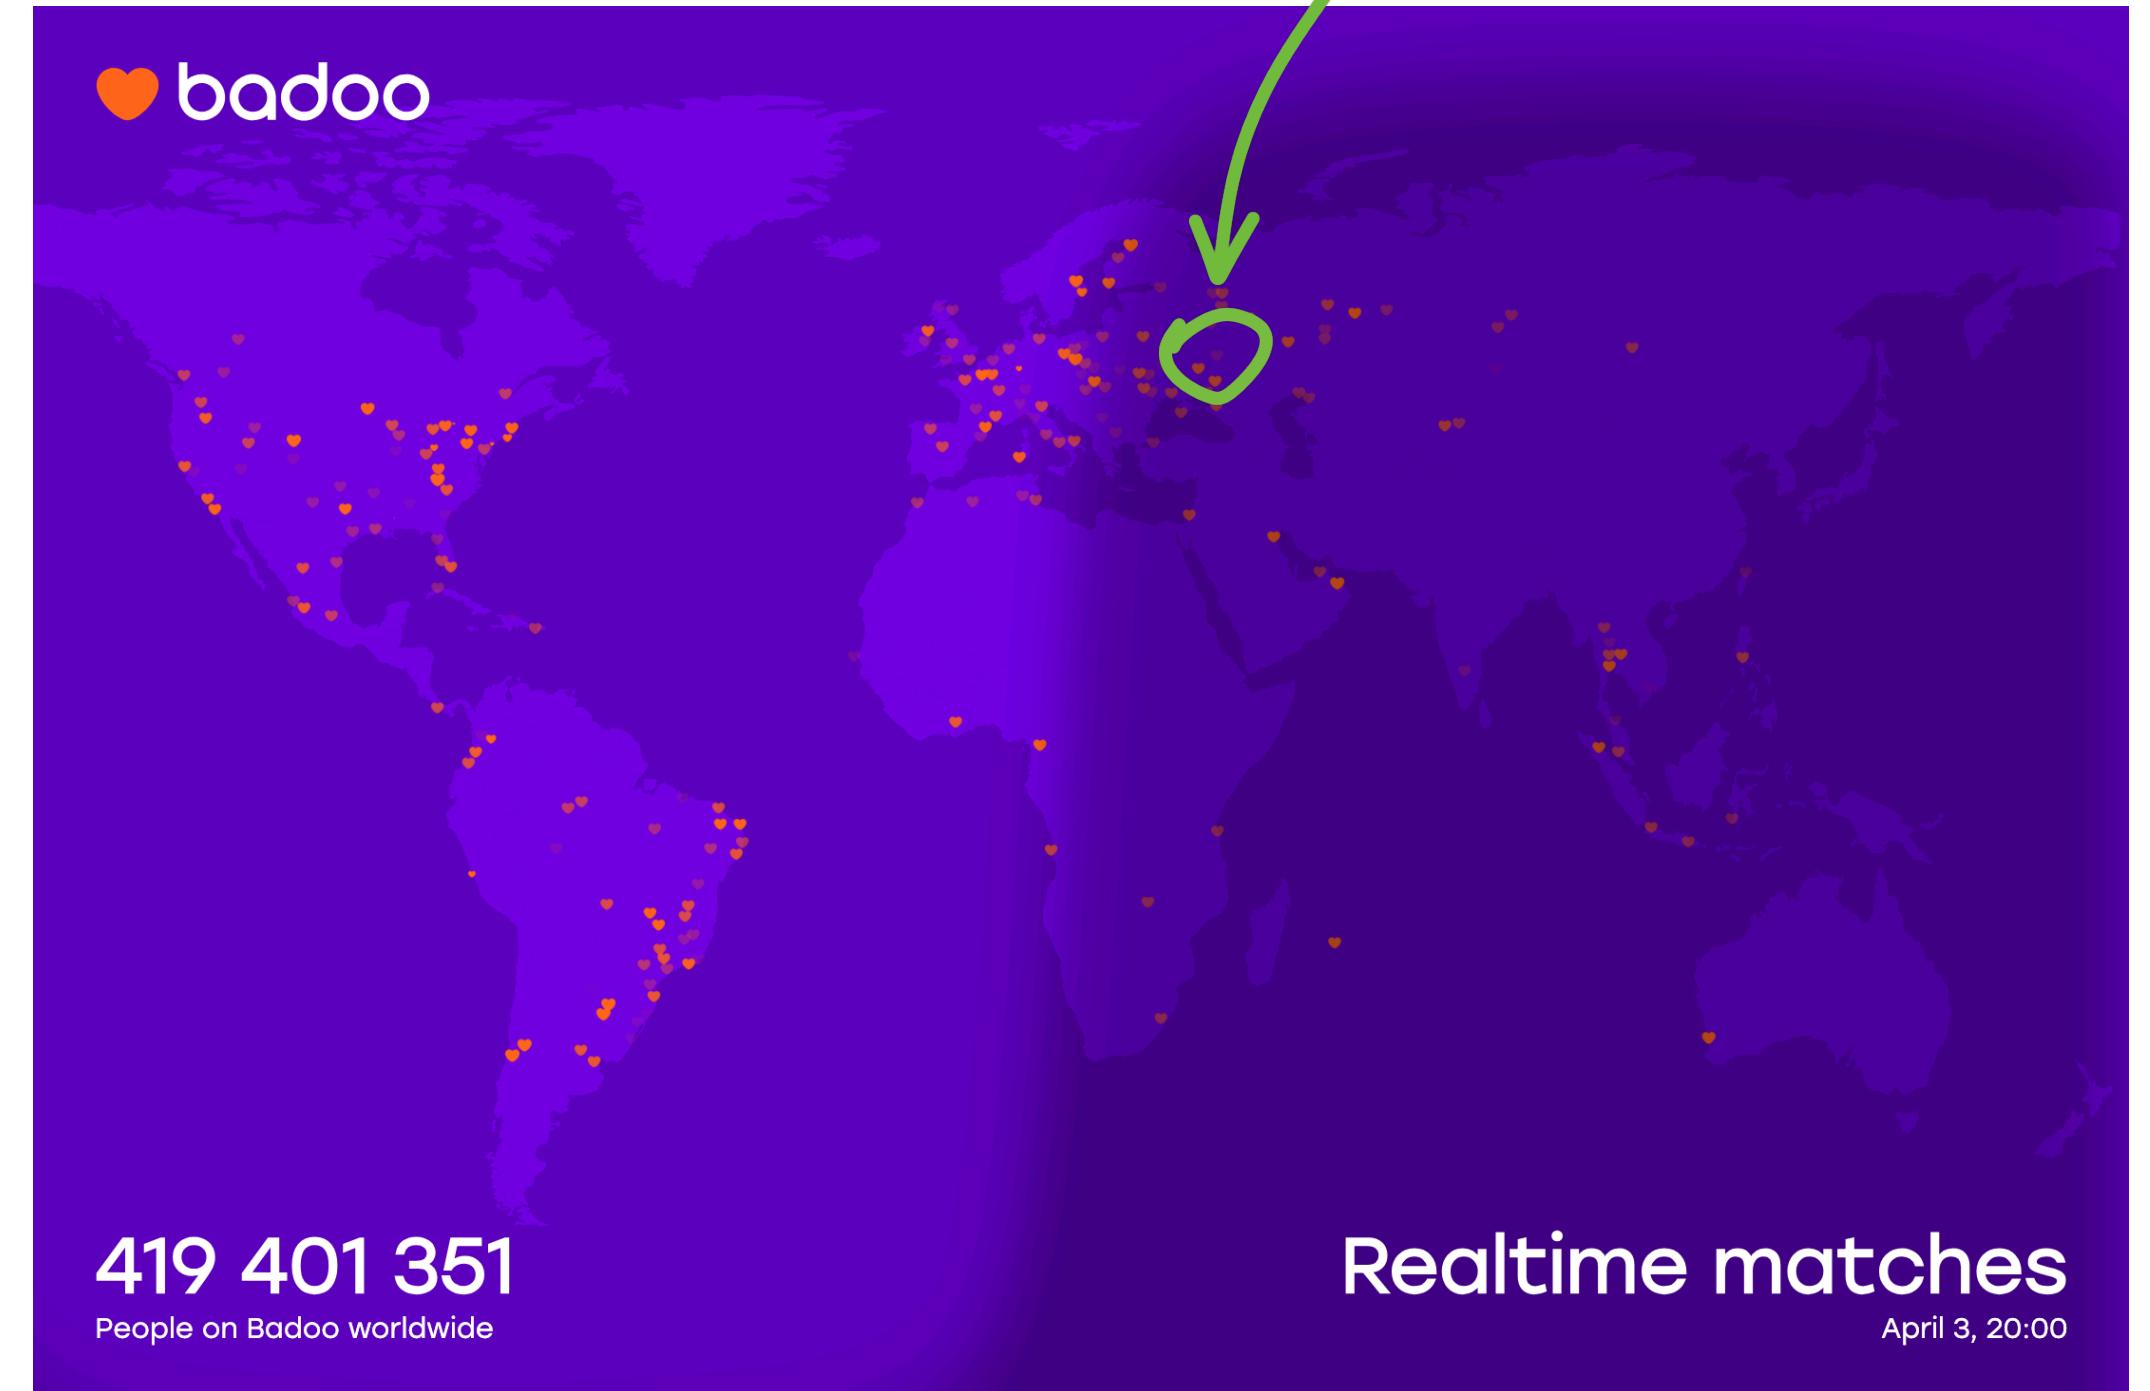
Why me?

Marko Kevac
marko@kevac.org



Badoo/Bumble

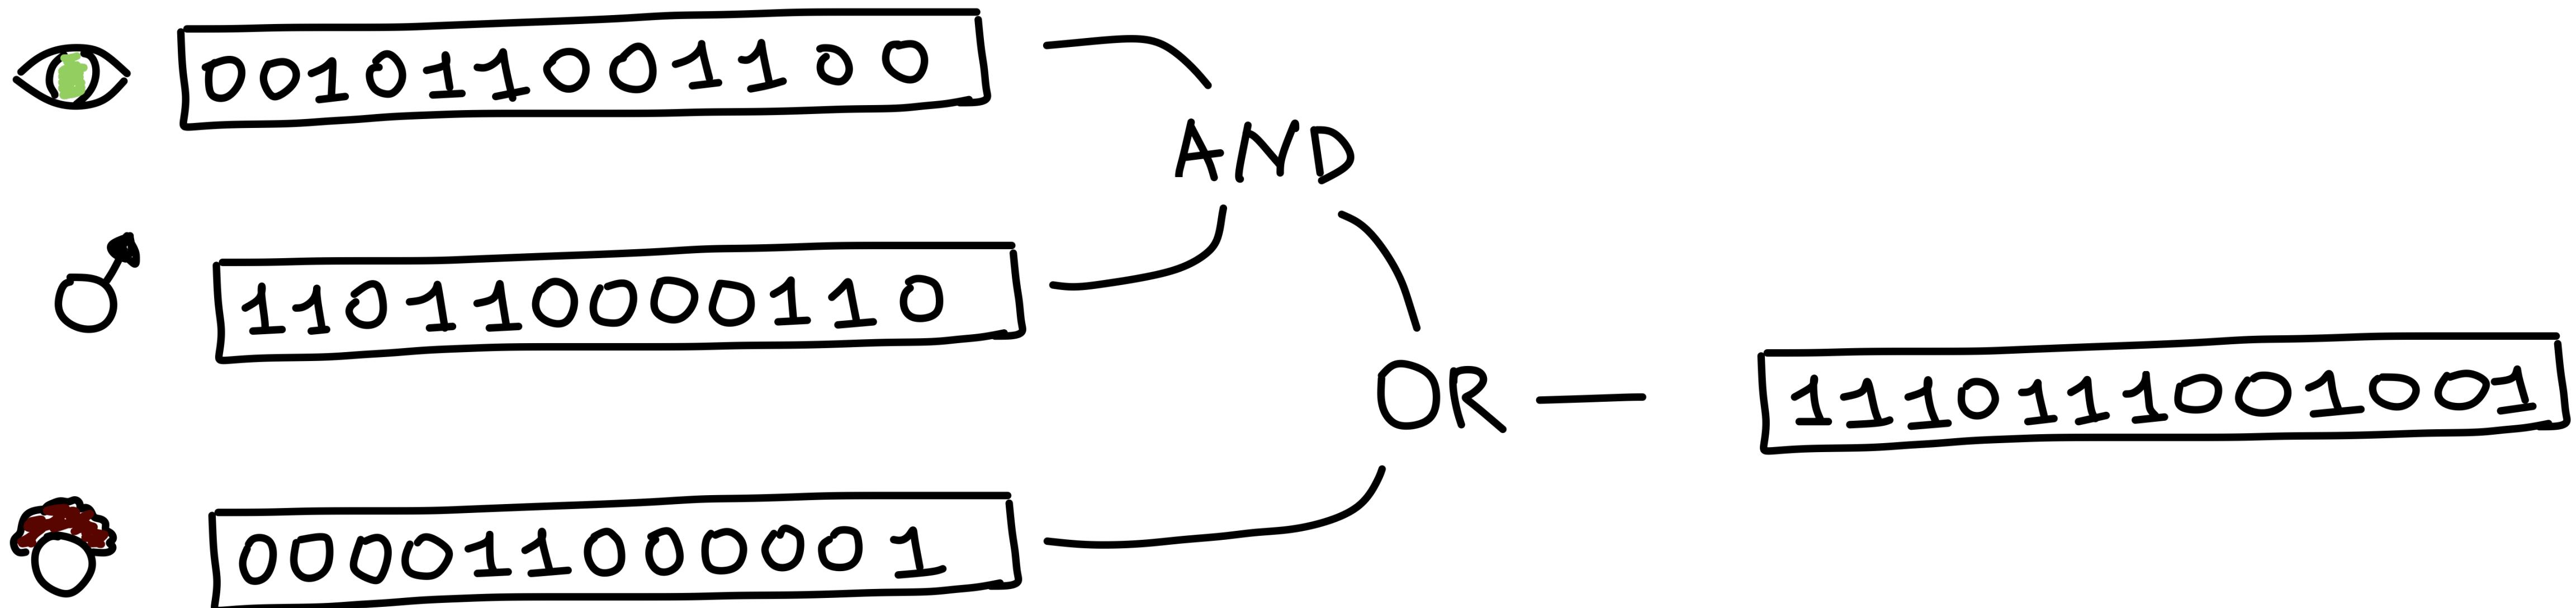
> 400 000 000 users worldwide



Searching for the best match for you using bitmap indexes since ~ 2010.

What is a Bitmap Index?

Bitmap / Bitset



Believed to be best for

Searches which combine several queries for different columns that have small cardinality.



“Give me results that have this
and this or this
and definitely not those...”

Small cardinality (few distinct values)

- Eye color
- Gender
- Marital status

Large cardinality

- Distance to the city center
- Number of likes received

Simplest example

Moscow restaurants

- Near Metro	0 . Pushkin
- Has private parking	1. Turandot
- Has terrace	2. Coffee Lovers
- Accepts reservations	3. Yunost
- Vegan friendly	4. Taras Bulba
- Expensive	...

Populating bitmaps

	0	1	2	3	4
near metro	1	0	0	1	1
PRIVATE PARK	1	0	0	0	0
TERRACE	0	1	0	0	1
RESERVATIONS	1	1	0	1	1
VEGAN	0	0	1	1	0
EXPENSIVE	1	1	0	0	0

Give me restaurants that are vegan friendly.

Give me restaurants that have terrace,
accept reservations, but are not expensive.

Give me restaurants that are vegan friendly.

Give me restaurants that have terrace,
accept reservations, but are not expensive.

Vegan friendly

	0	1	2	3	4
near metro	1	0	0	1	1
PRIVATE PARK	1	0	0	0	0
TERRACE	0	1	0	0	1
RESERVATIONS	1	1	0	1	1
VEGAN	0	0	1	1	0
EXPENSIVE	1	1	0	0	0

Give me restaurants that are vegan friendly.

Give me restaurants that have terrace,
accept reservations, but are not expensive.

More complex query

	0	1	2	3	4
near metro	1	0	0	1	1
PRIVATE PARK	1	0	0	0	0
TERRACE	0	1	0	0	1
RESERVATIONS	1	1	0	1	1
VEGAN	0	0	1	1	0
EXPENSIVE	1	1	0	0	0

Terrace AND Reservations

AND

NOT Expensive

More complex query

	0	1	2	3	4
near metro	1	0	0	1	1
PRIVATE PARK	1	0	0	0	0
TERRACE	0	1	0	0	1
RESERVATIONS	1	1	0	1	1
VEGAN	0	0	1	1	0
NOT EXPENS	0	0	1	1	1

Terrace AND Reservations

AND

NOT Expensive

More complex query

	0	1	2	3	4
near metro	1	0	0	1	1
PRIVATE PARK	1	0	0	0	0
TERRACE	0	1	0	0	1
RESERVATIONS	1	1	0	1	1
VEGAN	0	0	1	1	0
NOT EXPENS	0	0	1	1	1

Terrace AND Reservations

AND

NOT Expensive

Yunost

Bitmap index usage

Oracle DB



MySQL

PostgreSQL

Tarantool

Redis

MongoDB

ElasticSearch

Bitmap index usage

Oracle DB



Old fashioned

MySQL



Proposal

<https://dev.mysql.com/worklog/task/?id=1524>

PostgreSQL



Internal

Tarantool



Simple

Redis



<https://redis.io/commands/bitfield>

MongoDB



Proposal

<https://jira.mongodb.org/browse/SERVER-1723>

ElasticSearch



Internal

Bitmap index usage

Oracle DB	✓	Old fashioned
MySQL	—	Proposal
PostgreSQL	—	Internal
Tarantool	?	Simple
Redis	—	
MongoDB	—	Proposal
ElasticSearch	—	Internal
Pilosa	✓✓✓	Amazing

Implementation in Go

```
const restaurants = 65536
const bitmapLength = restaurants / 8 // 8192 bytes

var (
    nearMetro          = make([]byte, bitmapLength)
    privateParking     = make([]byte, bitmapLength)
    terrace            = make([]byte, bitmapLength)
    reservations       = make([]byte, bitmapLength)
    veganFriendly      = make([]byte, bitmapLength)
    expensive           = make([]byte, bitmapLength)
)
```

```
func fill(r *rand.Rand, b []byte, probability float32)
```

```
    fill(r, nearMetro, 0.1)
    fill(r, privateParking, 0.01)
    fill(r, terrace, 0.05)
    fill(r, reservations, 0.95)
    fill(r, veganFriendly, 0.2)
    fill(r, expensive, 0.1)
```

```
func indexes(a []byte) []int
```

Give me restaurants with a terrace that accept reservations,
but are not expensive.

terrace AND reservations AND (NOT expensive)

Give me restaurants with a terrace that accept reservations,
but are not expensive.

~~terrace AND reservations AND (NOT expensive)~~

terrace AND reservations **AND NOT** expensive

```
func and(a []byte, b []byte, res []byte) {
    for i := 0; i < len(a); i++ {
        res[i] = a[i] & b[i]
    }
}
```

```
func andnot(a []byte, b []byte, res []byte) {
    for i := 0; i < len(a); i++ {
        res[i] = a[i] & ^b[i]
    }
}
```

```
func BenchmarkSimpleBitmapIndex(b *testing.B) {
_, _, terrace, reservations, _, expensive := initData()

resBitmap := make([]byte, bitmapLength)
b.ResetTimer()

for i := 0; i < b.N; i++ {
    andnot(terrace, expensive, resBitmap)
    and(reservations, resBitmap, resBitmap)
}
}
```

name
SimpleBitmapIndex-12

time/op
10.8µs ± 0%

```
$ go test -run=None -bench=BenchmarkSimpleBitmapIndex$ -cpuprofile=cpu.out
```

```
$ pprof -http=:8080 ./simple.test ./cpu.out
```

```
.          .
.          .          for i := 0; i < b.N; i++ {
.          .          andnot(terrace, expensive, resBitmap)
.          680ms      .          MOVQ 0x20c8(SP), AX           simple_test.go:55
.          .          .          MOVQ AX, 0(SP)           simple_test.go:55
.          .          .          MOVQ 0x90(SP), CX           simple_test.go:55
.          .          .          MOVQ CX, 0x8(SP)           simple_test.go:55
.          .          .          MOVQ 0x98(SP), DX           simple_test.go:55
.          .          .          MOVQ DX, 0x10(SP)          simple_test.go:55
.          .          .          MOVQ 0x20d8(SP), BX           simple_test.go:55
.          .          .          MOVQ BX, 0x18(SP)           simple_test.go:55
.          .          .          MOVQ 0xb8(SP), SI           simple_test.go:55
.          .          .          MOVQ SI, 0x20(SP)           simple_test.go:55
.          .          .          MOVQ 0xc0(SP), DI           simple_test.go:55
.          .          .          MOVQ DI, 0x28(SP)           simple_test.go:55
.          .          .          LEAQ 0xc8(SP), R8           simple_test.go:55
.          .          .          MOVQ R8, 0x30(SP)           simple_test.go:55
.          .          .          MOVQ $0x2000, 0x38(SP)       simple_test.go:55
.          .          .          MOVQ $0x2000, 0x40(SP)       simple_test.go:55
.          680ms 10fd846:      .          CALL github.com/mkevac/gopherconrussia2019/simple.andnot(SB) simple_test.go:55
.
.          470ms      and(reservations, resBitmap, resBitmap)
.          .
.          .
}
```

```
$ go build -gcflags="-m -m"  
# github.com/mkevac/gopherconrussia2019/simple  
.simple.go:27:6: cannot inline and: unhandled op FOR  
.simple.go:71:6: cannot inline andnot: unhandled op FOR
```

cmd/compile: for-loops cannot be inlined #14768

<https://github.com/golang/go/issues/14768>

```
func andInlined(a []byte, b []byte, res []byte) {
    i := 0

    loop:
    if i < len(a) {
        res[i] = a[i] & b[i]
        i++
        goto loop
    }
}
```

```
$ go build -gcflags="-m -m" 2>&1 | grep andInlined
./simple.go:33:6: can inline andInlined as: func([]byte, []byte,
[]byte) { i := 0; loop: ; if i < len(a) { res[i] = a[i] & b[i]; i++;
goto loop } }
```

name

SimpleBitmapIndex-12

SimpleBitmapIndexInlined-12

time/op

10.8 μ s ± 0%

8.88 μ s ± 0%

~ 2 yrs

Bounds check elimination

```
/ \ . .
71   .   20ms          func andnot(a []byte, b []byte, res []byte) {
72   130ms 130ms          for i := 0; i < len(a); i++ {
73   450ms 450ms            res[i] = a[i] & ^b[i]
    90ms   90ms 10fcf9f:          MOVB R10, 0(BX)(AX*1)           simple.go:73
    :
    10ms   10ms 10fcfab:          MOVZX 0(R9)(AX*1), R10           simple.go:73
    70ms   70ms 10fcfb0:          CMPQ DI, AX                   simple.go:73
    .     . 10fcfb3:          JAE 0x10fcfdb             simple.go:73
120ms 120ms 10fcfb5:          MOVZX 0(R8)(AX*1), R11           simple.go:73
    30ms   30ms 10fcfba:          NOTL R11                  simple.go:73
    30ms   30ms 10fcfb0d:         ANDL R11, R10              simple.go:73
100ms 100ms 10fcfc0:          CMPQ SI, AX                 simple.go:73
    .     . 10fcfc3:          JB 0x10fcf9f               simple.go:73
    .     . 10fcfc5:          JMP 0x10fcfd1             simple.go:73
    .     . 10fcfc7:          MOVQ 0x10(SP), BP           simple.go:73
    .     . 10fcfc0c:         ADDQ $0x18, SP             simple.go:73
    .     . 10fcfd0:          RET                   simple.go:73
    .     . 10fcfd1:          MOVQ SI, CX                 simple.go:73
    .     . 10fcfd4:          CALL runtime.panicIndex(SB) simple.go:73
    .     . 10fcfd9:          UD2                   simple.go:73
    .     . 10fcfdb:          MOVQ DI, CX                 simple.go:73
    .     . 10fcfde:          CALL runtime.panicIndex(SB) simple.go:73
    .     . 10fcfe3:          UD2                   simple.go:73
74   .   .   } 
75   .   . }
```

```

func andNoBoundsCheck(a []byte, b []byte, res []byte) {
    if len(a) != len(b) || len(b) != len(res) {
        return
    }

    for i := 0; i < len(a); i++ {
        res[i] = a[i] & b[i]
    }
}

```

name	time/op
SimpleBitmapIndex-12	10.8µs ± 0%
SimpleBitmapIndexInlined-12	8.88µs ± 0%
SimpleBitmapIndexInlinedAndNoBoundsCheck-12	8.33µs ± 0%

~500ns

Bigger batches

```
const restaurants = 65536
const bitmapLength = restaurants / (8 * 8) // 8192 bytes (1024 elements)

nearMetro      = make([]uint64, bitmapLength)

func fill(r *rand.Rand, b []uint64, probability float32)

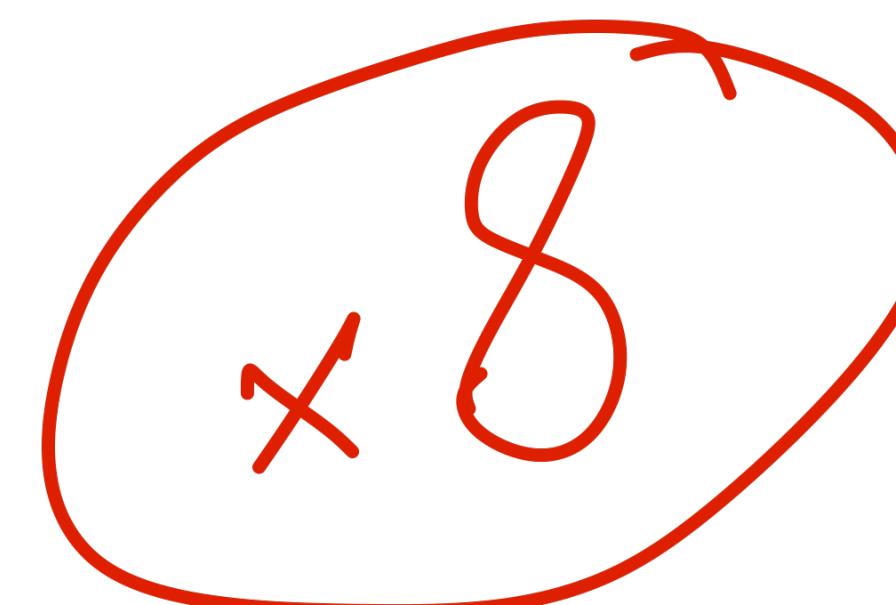
func indexes(a []uint64) []int

func andNoBoundsCheck(a []uint64, b []uint64, res []uint64) {
    if len(a) != len(b) || len(b) != len(res) {
        return
    }

    for i := 0; i < len(a); i++ {
        res[i] = a[i] & b[i]
    }
}
```

name	time/op
SimpleBitmapIndex-12	10.8µs ± 0%
SimpleBitmapIndexInlined-12	8.88µs ± 0%
SimpleBitmapIndexInlinedAndNoBoundsCheck-12	8.33µs ± 0%

name	time/op
BiggerBatchBitmapIndexNoBoundsCheck-12	1.06µs ± 0%

A red hand-drawn oval containing a large red 'X' and the number '8'.

SIMD

Single Instruction Multiple Data

- 16-byte chunks (SSE)
- 32-byte chunks (AVX, AVX2)
- 64-byte chunks (AVX512)

Vectorization

Go assembly

Not the real assembly, more like IRL.
Platform independent.

<https://www.youtube.com/watch?v=KINIAgRpkDA>

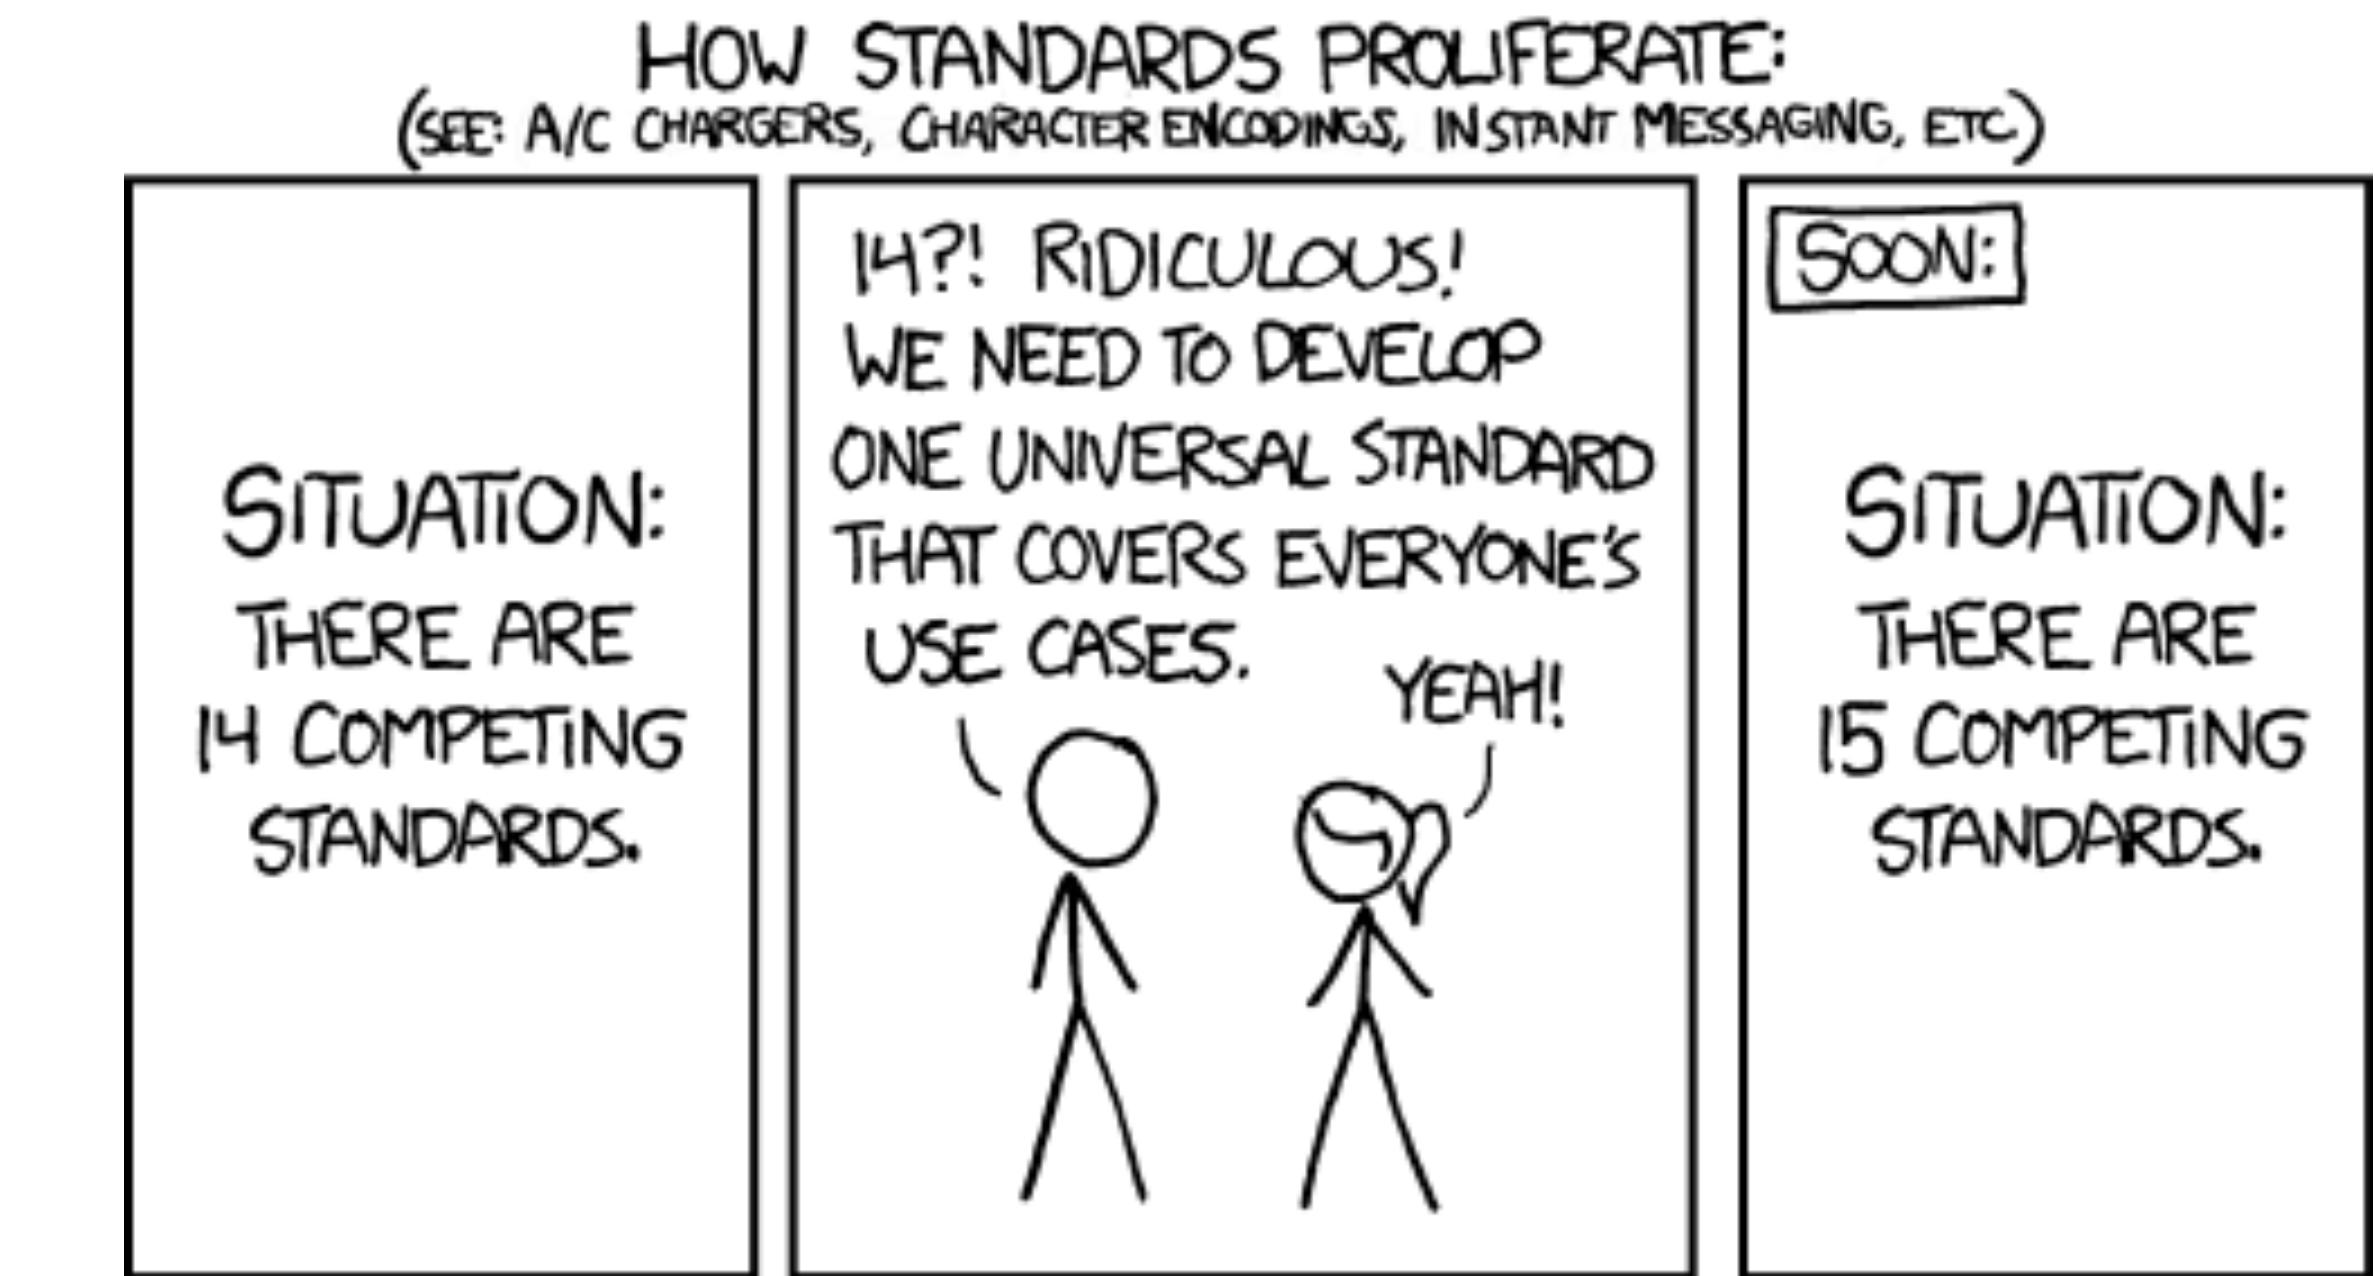
<https://habr.com/ru/company/badoo/blog/317864/> (rus)



Go assembly

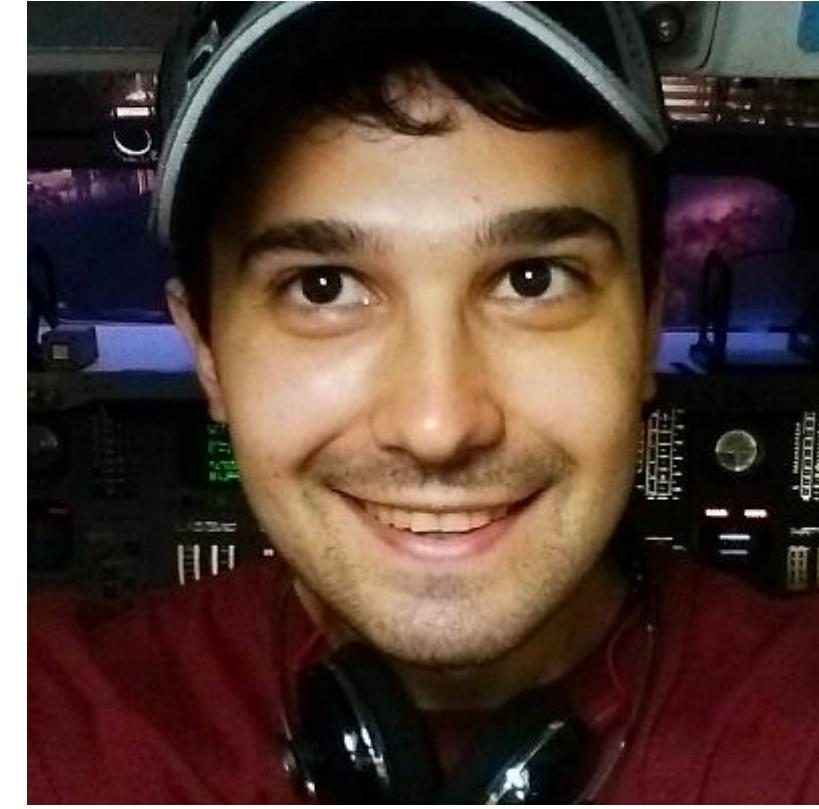
Not the real assembly, more like IRL.
Platform independent.

~~AT&T Intel Plan9!~~



<https://xkcd.com/927/>

No fun :-(



<https://github.com/Maratyszcza/PeachPy>



<https://github.com/mmcloughlin/avo>

Avo example

```
// +build ignore

package main

func main() {
    TEXT("Add", NOSPLIT, "func(x, y uint64) uint64")
    Doc("Add adds x and y.")
    x := Load(Param("x"), GP64())
    y := Load(Param("y"), GP64())
    ADDQ(x, y)
    Store(y, ReturnIndex(0))
    RET()
    Generate()
}
```

```
//go:generate go run asm.go -out add.s -stubs stub.go
```

Avo example

```
// +build ignore

package main

func main() {
    TEXT("Add", NOSPLIT, "func(x, y uint64) uint64")
    Doc("Add adds x and y.")
    x := Load(Param("x"), GP64())
    y := Load(Param("y"), GP64())
    ADDQ(x, y)
    Store(y, ReturnIndex(0))
    RET()
    Generate()
}
```

```
//go:generate go run asm.go -out add.s -stubs stub.go
```

src|S

```
// Code generated by command: go run asm.go -out add.s -stubs stub.go. DO NOT EDIT.

#include "textflag.h"

// func Add(x uint64, y uint64) uint64
TEXT ·Add(SB), NOSPLIT, $0-24
    MOVQ x+0(FP), AX
    MOVQ y+8(FP), CX
    ADDQ AX, CX
    MOVQ CX, ret+16(FP)
    RET
```

src|G
stub.go

```
// Code generated by command: go run asm.go -out add.s -stubs stub.go. DO NOT EDIT.

package add

// Add adds x and y.
func Add(x uint64, y uint64) uint64
```

Scalar version

```
TEXT("andnotScalarFaster", NOSPLIT, "func(a []byte, b []byte, res []byte)")

a := Mem{Base: Load(Param("a").Base(), GP64())}
b := Mem{Base: Load(Param("b").Base(), GP64())}
res := Mem{Base: Load(Param("res").Base(), GP64())}
n := Load(Param("a").Len(), GP64())

ir := GP64()
XORQ(ir, ir)

ar := GP64()
br := GP64()

Label("loop")
CMPQ(n, ir)
JE(LabelRef("done"))

MOVQ(a.Idx(ir, 1), ar)
MOVQ(b.Idx(ir, 1), br)

ANDNQ(ar, br, br)

MOVQ(br, res.Idx(ir, 1))

ADDQ(Imm(8), ir)

JMP(LabelRef("loop"))

Label("done")
RET()
Generate()
```

The diagram illustrates the flow of control in the scalar version of the function. Red arrows connect the initial memory loads (a, b, res) to their respective assembly counterparts (MOVQ). Another red arrow connects the 'done' label back to the 'Generate()' function at the bottom.

Scalar version

```
TEXT("andnotScalarFaster", NOSPLIT, "func(a []byte, b []byte, res []byte)")

a := Mem{Base: Load(Param("a").Base(), GP64())}
b := Mem{Base: Load(Param("b").Base(), GP64())}
res := Mem{Base: Load(Param("res").Base(), GP64())}
n := Load(Param("a").Len(), GP64())

ir := GP64()
XORQ(ir, ir)

ar := GP64()
br := GP64()

Label("loop")
CMPQ(n, ir)
JE(LabelRef("done"))

MOVQ(a.Idx(ir, 1), ar)
MOVQ(b.Idx(ir, 1), br)

ANDNQ(ar, br, br)

MOVQ(br, res.Idx(ir, 1))

ADDQ(Imm(8), ir)

JMP(LabelRef("loop"))

Label("done")
RET()
Generate()
```

Scalar version

```
TEXT("andnotScalarFaster", NOSPLIT, "func(a []byte, b []byte, res []byte)")
```

```
a := Mem{Base: Load(Param("a").Base(), GP64())}  
b := Mem{Base: Load(Param("b").Base(), GP64())}  
res := Mem{Base: Load(Param("res").Base(), GP64())}  
n := Load(Param("a").Len(), GP64())
```

```
ir := GP64()  
XORQ(ir, ir)
```

```
ar := GP64()  
br := GP64()
```

```
Label("loop")  
CMPQ(n, ir)  
JE(LabelRef("done"))
```

```
MOVQ(a.Idx(ir, 1), ar)  
MOVQ(b.Idx(ir, 1), br)
```

```
ANDNQ(ar, br, br)
```

```
MOVQ(br, res.Idx(ir, 1))
```

```
ADDQ(Imm(8), ir)
```

```
JMP(LabelRef("loop"))
```

```
Label("done")  
RET()  
Generate()
```

```
// func andnotScalarFaster(a []byte, b []byte, res []byte)
TEXT ·andnotScalarFaster(SB), NOSPLIT, $0-72
    MOVQ a_base+0(FP), AX
    MOVQ b_base+24(FP), CX
    MOVQ res_base+48(FP), DX
    MOVQ a_len+8(FP), BX
    XORQ BP, BP
```

loop:

```
CMPQ BX, BP
JE done
MOVQ (AX)(BP*1), SI
MOVQ (CX)(BP*1), DI
ANDNQ SI, DI, DI
MOVQ DI, (DX)(BP*1)
ADDQ $0x08, BP
JMP loop
```

done:

```
RET
```

name	time/op
SimpleBitmapIndex-12	$10.8\mu\text{s} \pm 0\%$
SimpleBitmapIndexInlined-12	$8.88\mu\text{s} \pm 0\%$
SimpleBitmapIndexInlinedAndNoBoundsCheck-12	$8.33\mu\text{s} \pm 0\%$

name	time/op
BiggerBatchBitmapIndexNoBoundsCheck-12	$1.06\mu\text{s} \pm 0\%$

name	time/op
SimpleScalarFasterBitmapIndex-8	$1.04\mu\text{s} \pm 0\%$

Inlining assembly code is not possible: <https://github.com/golang/go/issues/26891>

Vector (SIMD) version

```
var unroll = 8

TEXT("andnotSIMD", NOSPLIT, "func(a []byte, b []byte, res []byte)")

a := Mem{Base: Load(Param("a").Base(), GP64())}
b := Mem{Base: Load(Param("b").Base(), GP64())}
res := Mem{Base: Load(Param("res").Base(), GP64())}
n := Load(Param("a").Len(), GR64())

blocksize := 32 * unroll

Label("blockloop")
CMPQ(n, U32(blocksize))
JL(LabelRef("tail"))
```

```
// Load b.  
bs := make([]VecVirtual, unroll)  
for i := 0; i < unroll; i++ {  
    bs[i] = YMM()  
}  
  
for i := 0; i < unroll; i++ {  
    VMOVUPS(b.Offset(32*i), bs[i])  
}  
  
// The actual operation.  
for i := 0; i < unroll; i++ {  
    VPANDN(a.Offset(32*i), bs[i], bs[i])  
}  
  
for i := 0; i < unroll; i++ {  
    VMOVUPS(bs[i], res.Offset(32*i))  
}
```

8

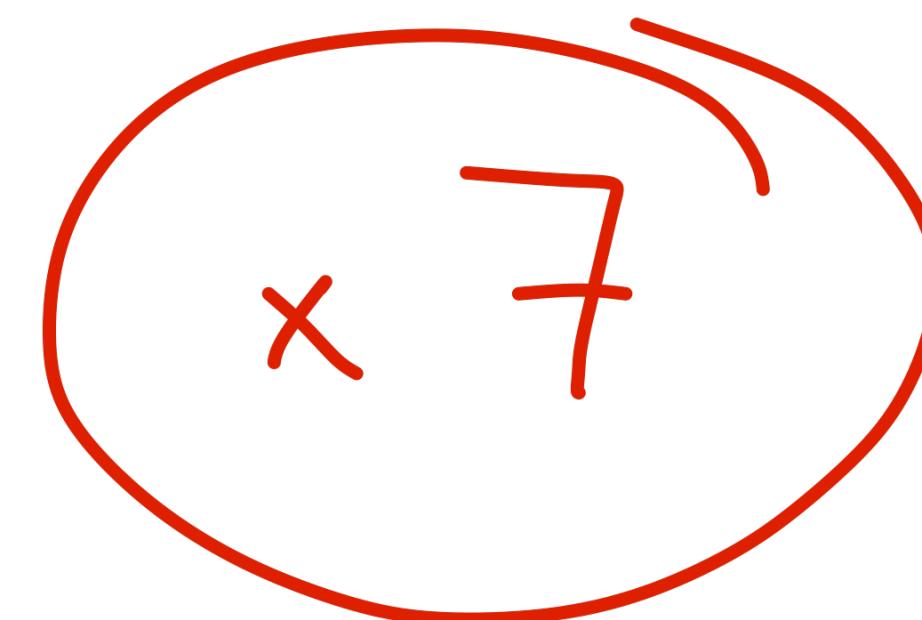
ADDQ(U32(blocksize), a.Base)
ADDQ(U32(blocksize), b.Base)
ADDQ(U32(blocksize), res.Base)
SUBQ(U32(blocksize), n)
JMP(LabelRef("blockloop"))

Label("tail")
RET()

name	time/op
SimpleBitmapIndex-12	$10.8\mu\text{s} \pm 0\%$
SimpleBitmapIndexInlined-12	$8.88\mu\text{s} \pm 0\%$
SimpleBitmapIndexInlinedAndNoBoundsCheck-12	$8.33\mu\text{s} \pm 0\%$

name	time/op
BiggerBatchBitmapIndexNoBoundsCheck-12	$1.06\mu\text{s} \pm 0\%$

name	time/op
SimpleSIMDBitmapIndex-8	$154\text{ns} \pm 0\%$
SimpleScalarFasterBitmapIndex-8	$1.04\mu\text{s} \pm 0\%$



name	time/op
SimpleBitmapIndex-12	$10.8\mu\text{s} \pm 0\%$
SimpleBitmapIndexInlined-12	$8.88\mu\text{s} \pm 0\%$
SimpleBitmapIndexInlinedAndNoBoundsCheck-12	$8.33\mu\text{s} \pm 0\%$

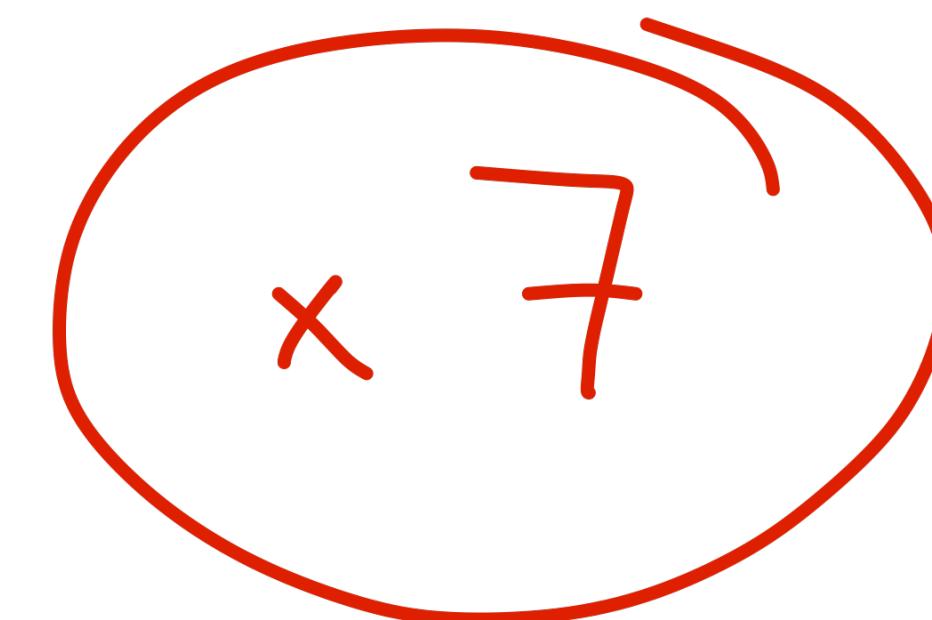
name	time/op
BiggerBatchBitmapIndexNoBoundsCheck-12	$1.06\mu\text{s} \pm 0\%$

name	time/op
SimpleSIMDBitmapIndex-8	$154\text{ns} \pm 0\%$
SimpleScalarFasterBitmapIndex-8	$1.04\mu\text{s} \pm 0\%$

AVX 512

PREFETCHING

JIT



Bitmap Index “Problems”

1. High cardinality problem
2. High throughput problem
3. Non-trivial queries

High cardinality problem

Representation leads to sparsely populated bitmaps.

0000000000001000000000000000000000
000000000000000000000000000000000000
000000000000000000000000000000000000



Different representation.

Bit per value

0001000000

cardinality 10
value 3

Binary representation

0011 4 bit width

Compression.

PWAH
EWAH

00000000001111
10 4

Run-length encoding

01014

Roaring bitmaps.

Bitmaps

Arrays

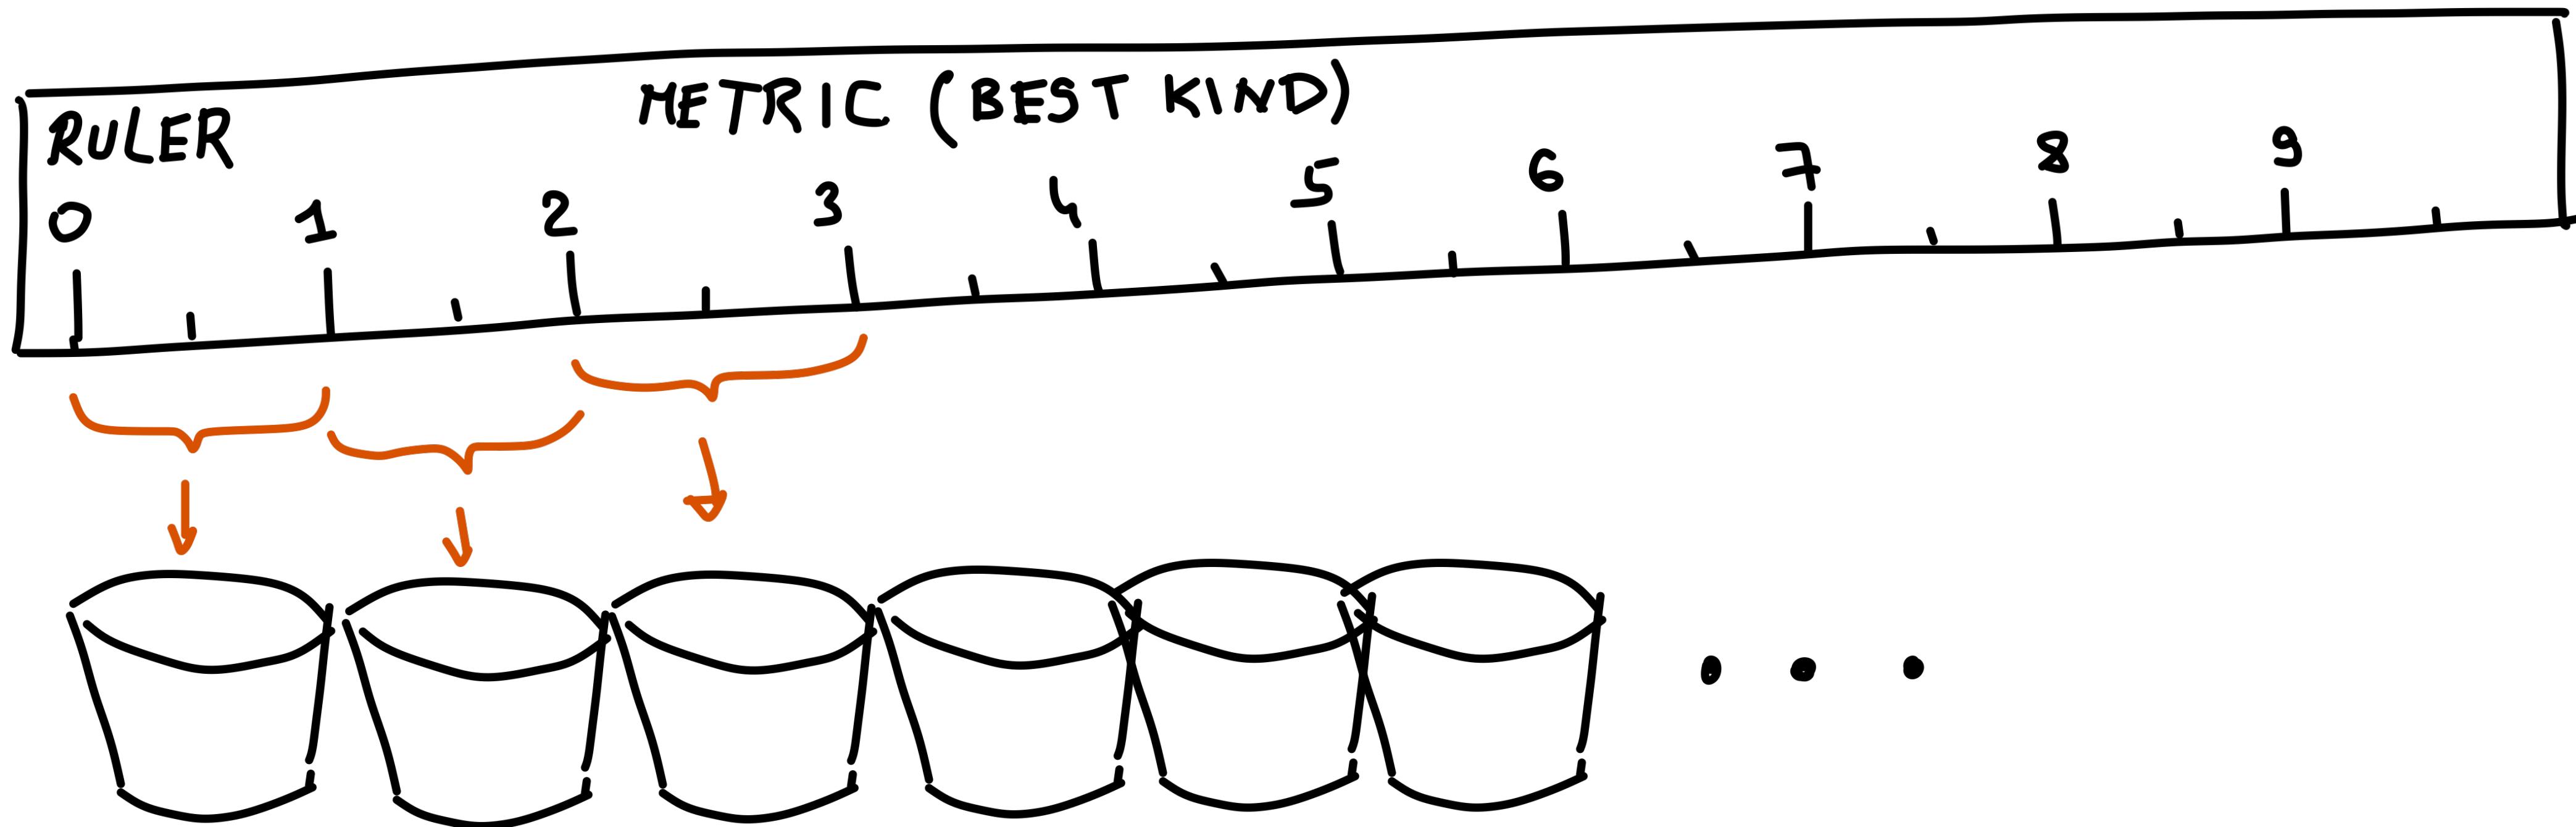
Bit runs

<https://www.roaringbitmap.org/>

<https://arxiv.org/pdf/1603.06549.pdf>

Apache Lucene
Apache Spark
InfluxDB
Netflix Atlas
Bleve

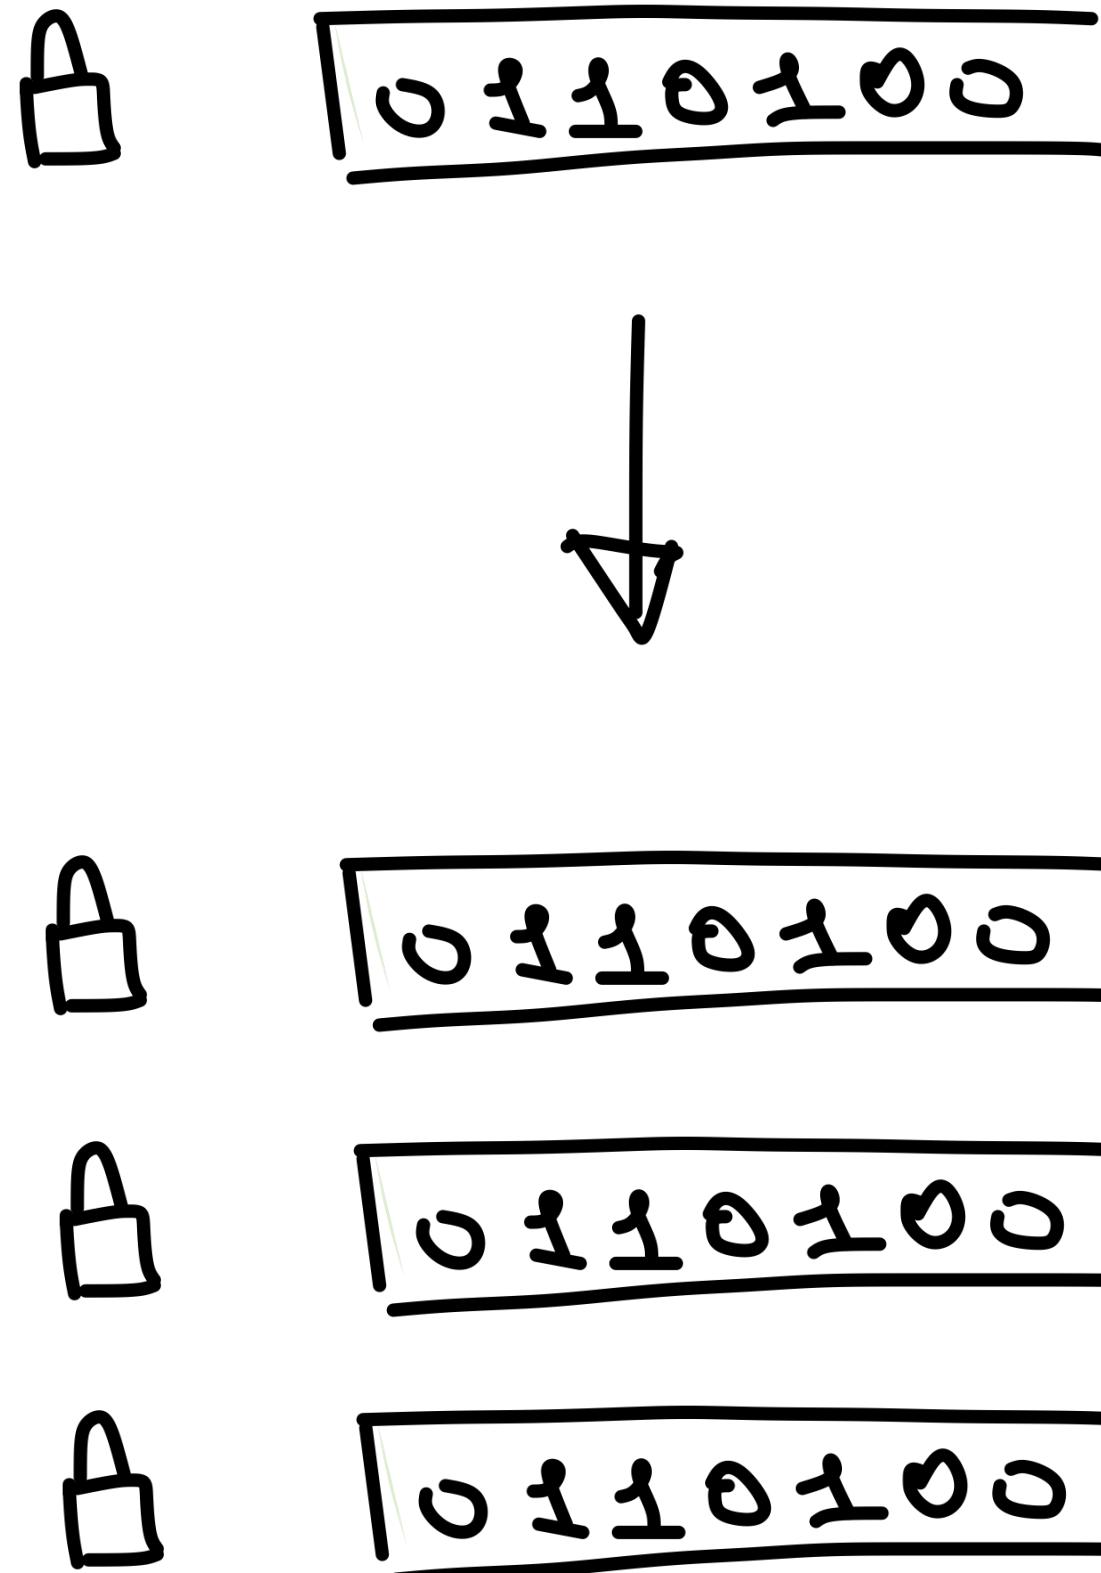
Binning.



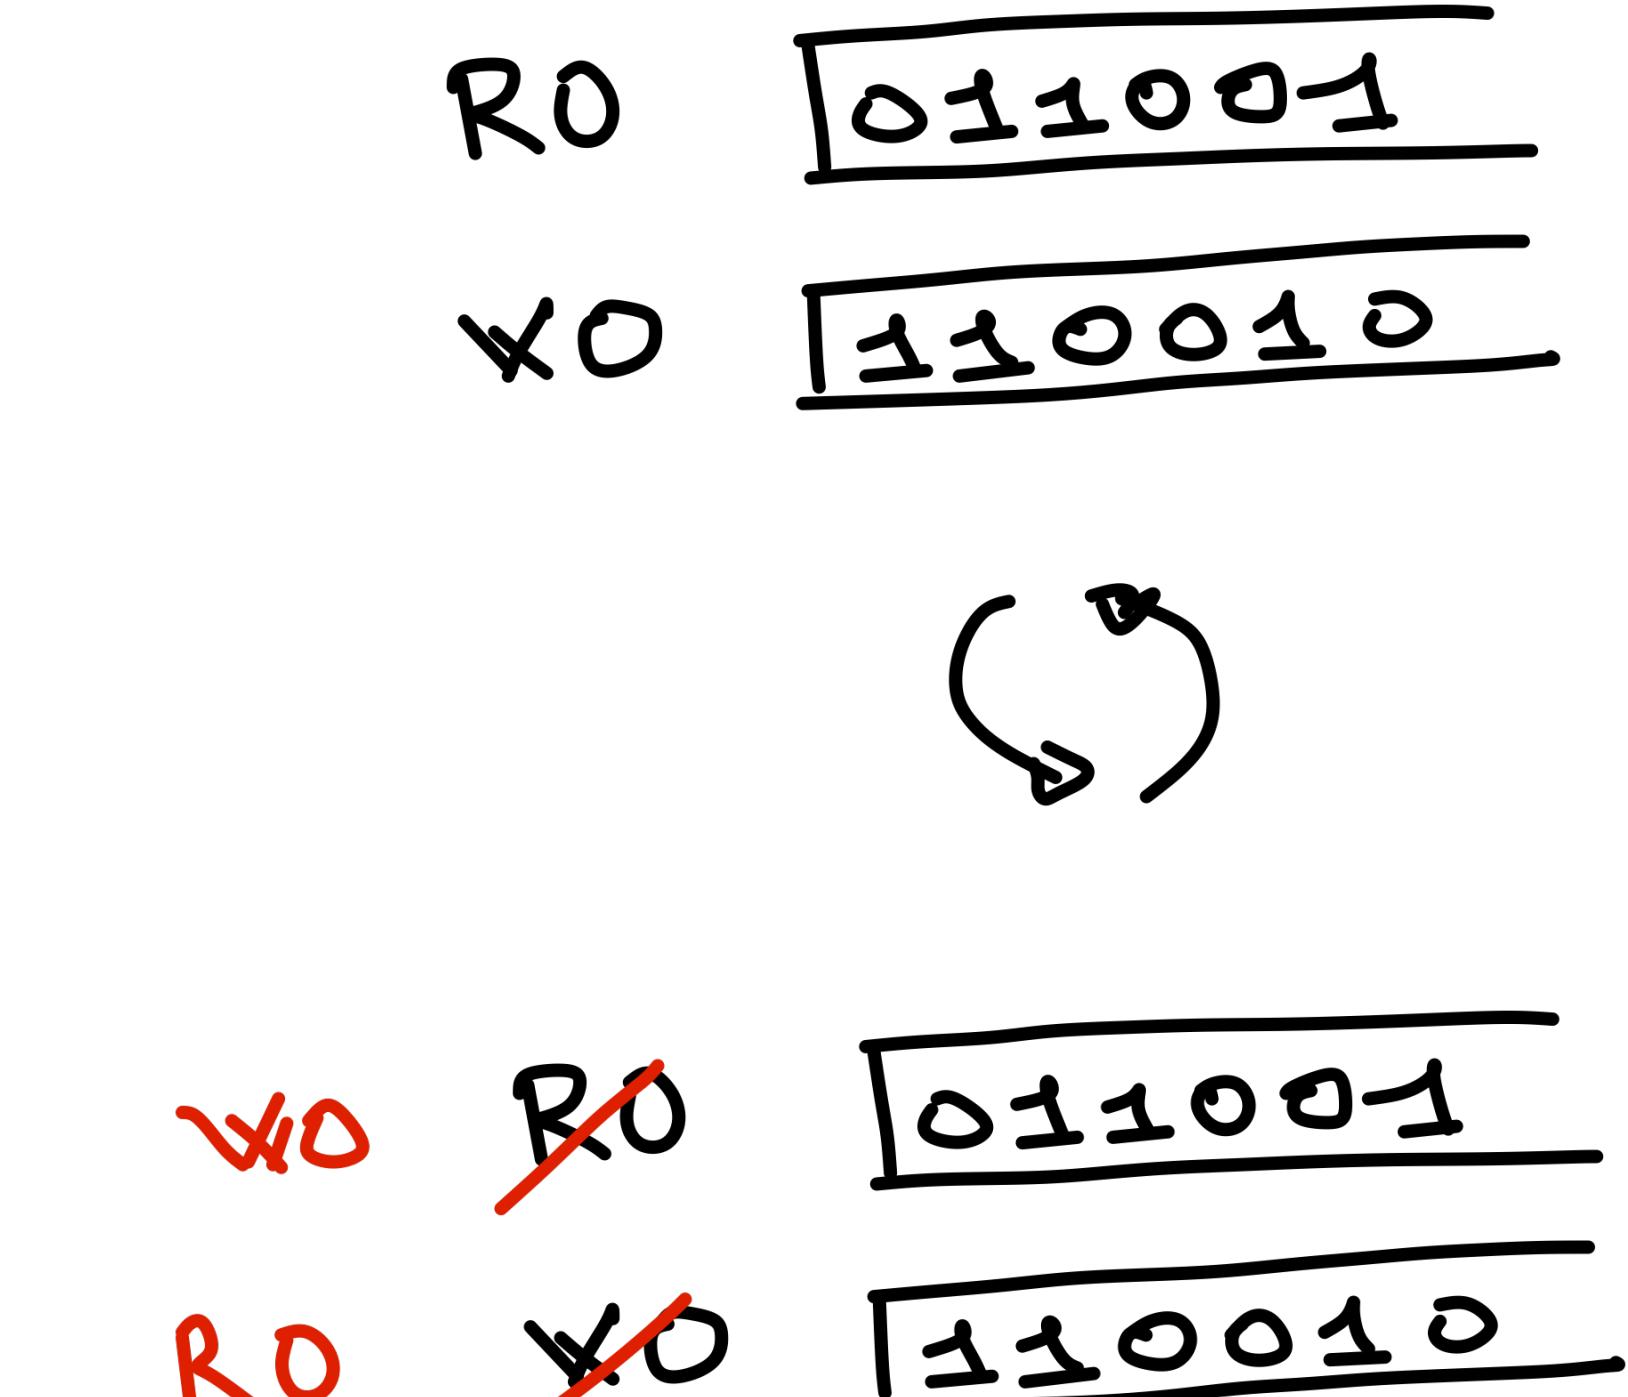
Height (float, (0, inf), inf cardinality) -> Binned height (uint, (50, 250), cardinality 200)

High throughput problem

Sharding



Versioning



Non-trivial queries: range queries

Straightforward solution

“Give me hotel rooms that cost from 200 to 300 dollars a night.”

“Give me hotel rooms that cost 200 dollars a night.”

OR

“Give me hotel rooms that cost 201 dollars a night.”

OR

“Give me hotel rooms that cost 202 dollars a night.”

OR

“Give me hotel rooms that cost 203 dollars a night.”

...

OR

“Give me hotel rooms that cost 300 dollars a night.”

Non-trivial queries: range queries

Binning

“Give me hotel rooms that cost from 200 to 300 dollars a night.”

“Give me hotel rooms that cost 200-250 dollars a night.”

OR

“Give me hotel rooms that cost 250-300 dollars a night.”

x 50

Non-trivial queries: range queries

Range-encoded bitmaps

“Give me hotel rooms that cost from 200 to 300 dollars a night.”

200 300

0000011111111111

“Give me hotel rooms that cost ≤ 300 dollars a night.”

AND
NOT

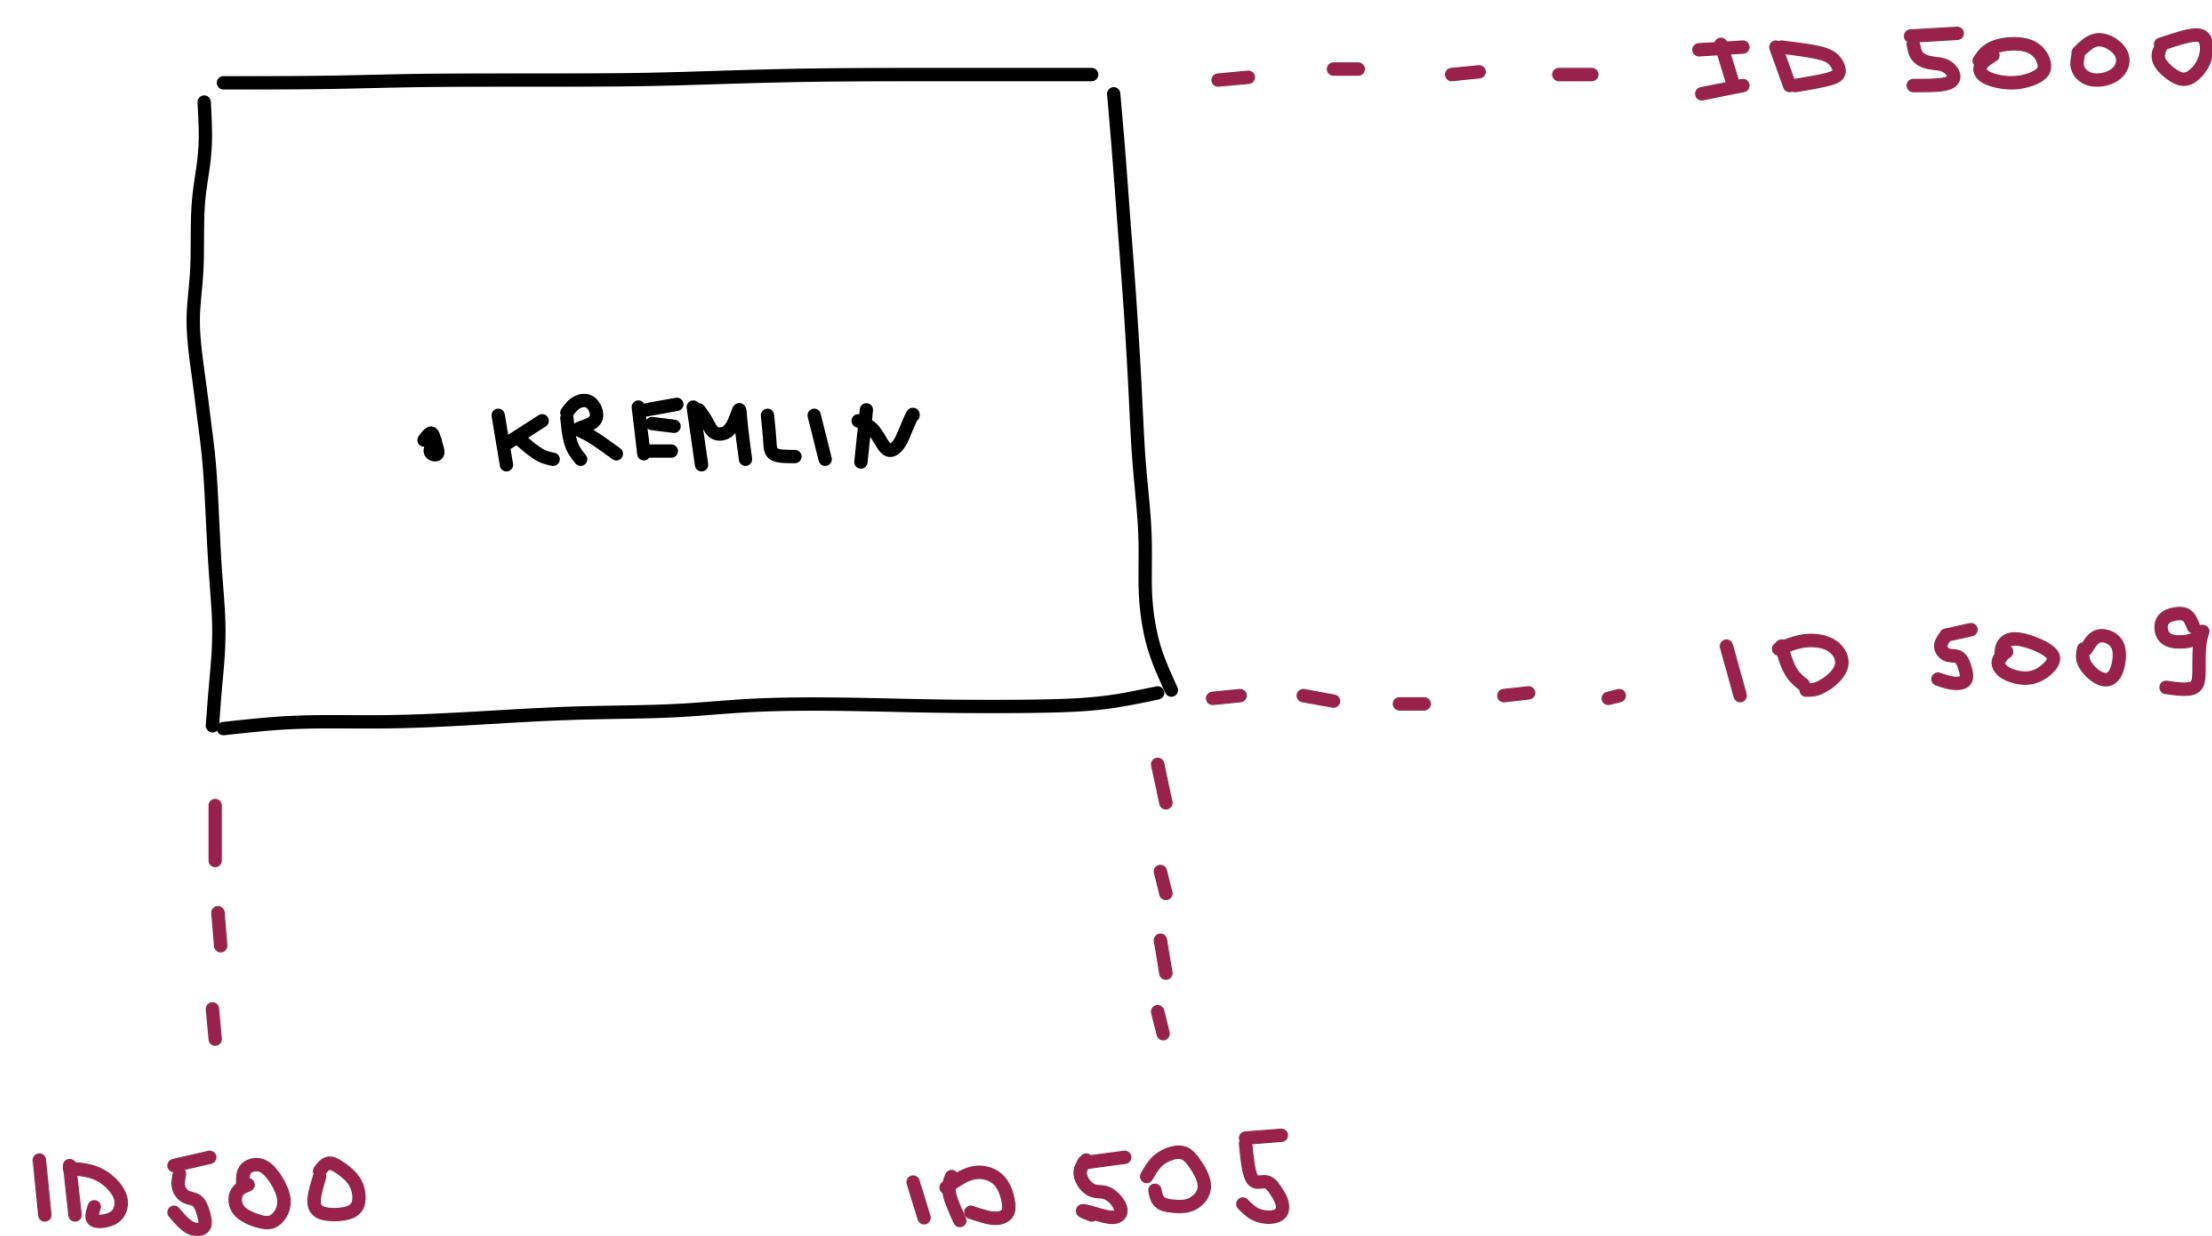
00000000000000111

“Give me hotel rooms that cost ≤ 199 dollars a night.”

Non-trivial queries: geo queries

Google S2 ([Official site](#), [Gophercon Russia 2018 talk](#))

Badoo Geotri ([April 1 article on Habr](#))



WHERE ID VERTICAL
IN (5000, 5009) AND
ID HORIZONTAL
IN (500, 505)

Ready solutions?

1. Roaring bitmaps
2. Pilosa
3. Maybe mainstream DBMS with get bitmap index support?

Ready solutions?

1. Roaring bitmaps

2. Pilosa

3. Maybe mainstream DBMS with get bitmap index support?

CONTAINER

BITWISE OPERATIONS

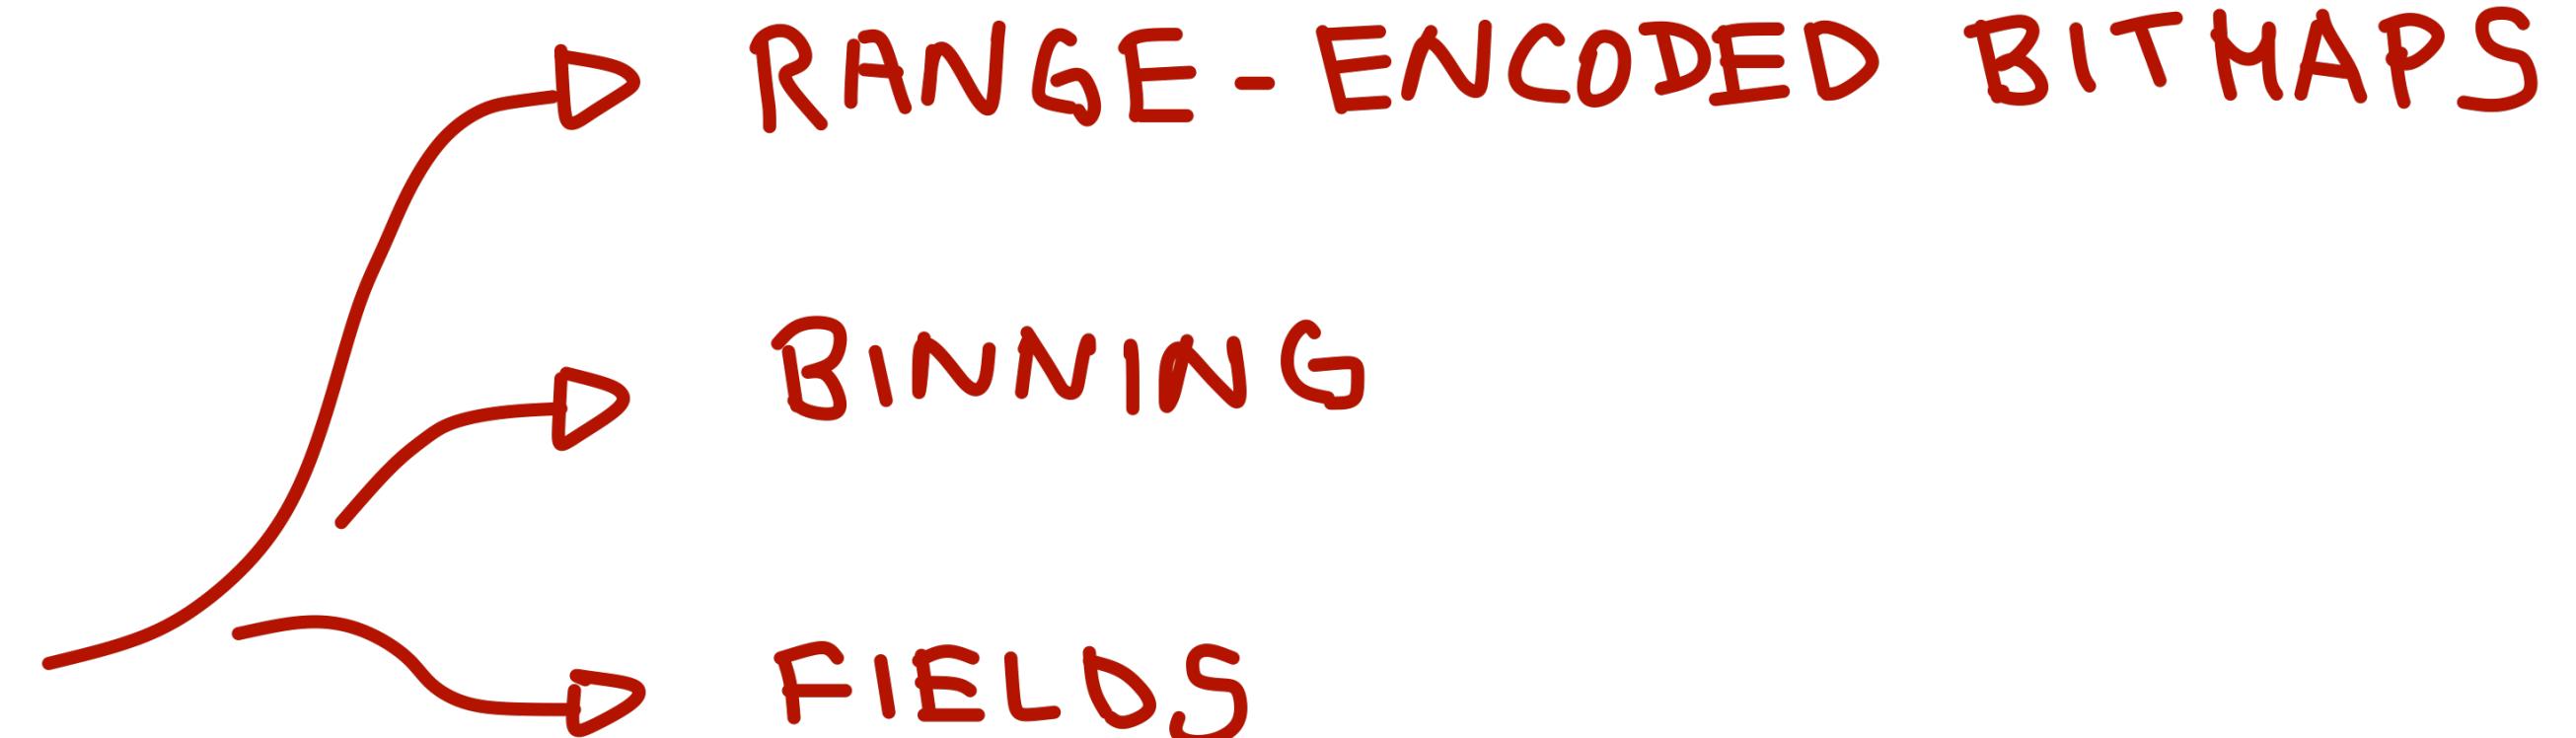
NO SIMD

Ready solutions?

1. Roaring bitmaps

2. Pilosa

3. Maybe mainstream DBMS with get bitmap index support?



Pilosa example

```
client = pilosa.DefaultClient()

// Retrieve the schema
schema, err = client.Schema()

// Create an Index object
myindex = schema.Index("myindex")

terraceField = myindex.Field("terrace")
reservationsField = myindex.Field("reservations")
expensiveField = myindex.Field("expensive")

// make sure the index and the field exists on the server
err = client.SyncSchema(schema)

fill(r, terraceField, 0.05)
fill(r, reservationsField, 0.95)
fill(r, expensiveField, 0.1)

response, err = client.Query(myindex.Intersect(
    terraceField.Row(0),
    myindex.Not(expensiveField.Row(0)),
    reservationsField.Row(0)))
```

Pilosa example

```
client = pilosa.DefaultClient()

// Retrieve the schema
schema, err = client.Schema()

// Create an Index object
myindex = schema.Index("myindex")

terraceField = myindex.Field("terrace")
reservationsField = myindex.Field("reservations")
expensiveField = myindex.Field("expensive")

// make sure the index and the field exists on the server
err = client.SyncSchema(schema)

fill(r, terraceField, 0.05)
fill(r, reservationsField, 0.95)
fill(r, expensiveField, 0.1)

response, err = client.Query(myindex.Intersect(
    terraceField.Row(0),
    myindex.Not(expensiveField.Row(0)),
    reservationsField.Row(0)))
```

Pilosa example

```
client = pilosa.DefaultClient()

// Retrieve the schema
schema, err = client.Schema()

// Create an Index object
myindex = schema.Index("myindex")

terraceField = myindex.Field("terrace")
reservationsField = myindex.Field("reservations")
expensiveField = myindex.Field("expensive")

// make sure the index and the field exists on the server
err = client.SyncSchema(schema)

fill(r, terraceField, 0.05)
fill(r, reservationsField, 0.95)
fill(r, expensiveField, 0.1)

response, err = client.Query(myindex.Intersect(
    terraceField.Row(0),
    myindex.Not(expensiveField.Row(0)),
    reservationsField.Row(0)))
```

Pilosa example

```
client = pilosa.DefaultClient()

// Retrieve the schema
schema, err = client.Schema()

// Create an Index object
myindex = schema.Index("myindex")

terraceField = myindex.Field("terrace")
reservationsField = myindex.Field("reservations")
expensiveField = myindex.Field("expensive")

// make sure the index and the field exists on the server
err = client.SyncSchema(schema)

fill(r, terraceField, 0.05)
fill(r, reservationsField, 0.95)
fill(r, expensiveField, 0.1)

response, err = client.Query(myindex.Intersect(
    terraceField.Row(0),
    myindex.Not(expensiveField.Row(0)),
    reservationsField.Row(0)))
```

Pilosa example

```
client = pilosa.DefaultClient()

// Retrieve the schema
schema, err = client.Schema()

// Create an Index object
myindex = schema.Index("myindex")

terraceField = myindex.Field("terrace")
reservationsField = myindex.Field("reservations")
expensiveField = myindex.Field("expensive")

// make sure the index and the field exists on the server
err = client.SyncSchema(schema)

fill(r, terraceField, 0.05)
fill(r, reservationsField, 0.95)
fill(r, expensiveField, 0.1)

response, err = client.Query(myindex.Intersect(
    terraceField.Row(0),
    myindex.Not(expensiveField.Row(0)),
    reservationsField.Row(0)))
```

Pilosa example

```
client = pilosa.DefaultClient()

// Retrieve the schema
schema, err = client.Schema()

// Create an Index object
myindex = schema.Index("myindex")

terraceField = myindex.Field("terrace")
reservationsField = myindex.Field("reservations")
expensiveField = myindex.Field("expensive")

// make sure the index and the field exists on the server
err = client.SyncSchema(schema)

fill(r, terraceField, 0.05)
fill(r, reservationsField, 0.95)
fill(r, expensiveField, 0.1)

response, err = client.Query(myindex.Intersect(
    terraceField.Row(0),
    myindex.Not(expensiveField.Row(0)),
    reservationsField.Row(0)))
```

```
$ go run pilosa.go
2019/03/30 20:41:12 filling the data...
2019/03/30 20:41:47 finished filling the data
2019/03/30 20:41:47 got 2796 columns
```

Ready solutions?

1. Roaring bitmaps
2. Pilosa
3. Maybe mainstream DBMS will get bitmap index support?

Closing words

- Bitmap indexes
- Performance optimizations

Thank you

<http://bit.ly/bitmapindexes>

<https://github.com/mkevac/gopherconrussia2019>